Data Mining

CS4821-CS5831-s24

Laura E. Brown

Jan. 10, 2024

Some slides from G. Piatetsky-Shapiro; Han, Kamber, & Pei; M. Hahsler
P. Smyth; C. Volinsky; Tan, Steinbach, & Kumar; J. Taylor; G. Dong; L. Hannah

Today's Agenda

- Examples of Data Mining (continue from 01.intro.part1)
- Multidimensional View of Data Mining
- Tools for Data Mining
- Data Mining vs. Privacy

Multi-dimensional View of Data Mining

- What kinds of **data** can be mined?
- What kinds of **patterns** can be mined?
- What kinds of **techniques** are used?
- What kinds of **applications** are targeted?

What kind of data?

- Spreadsheets
- Flat file, vector data

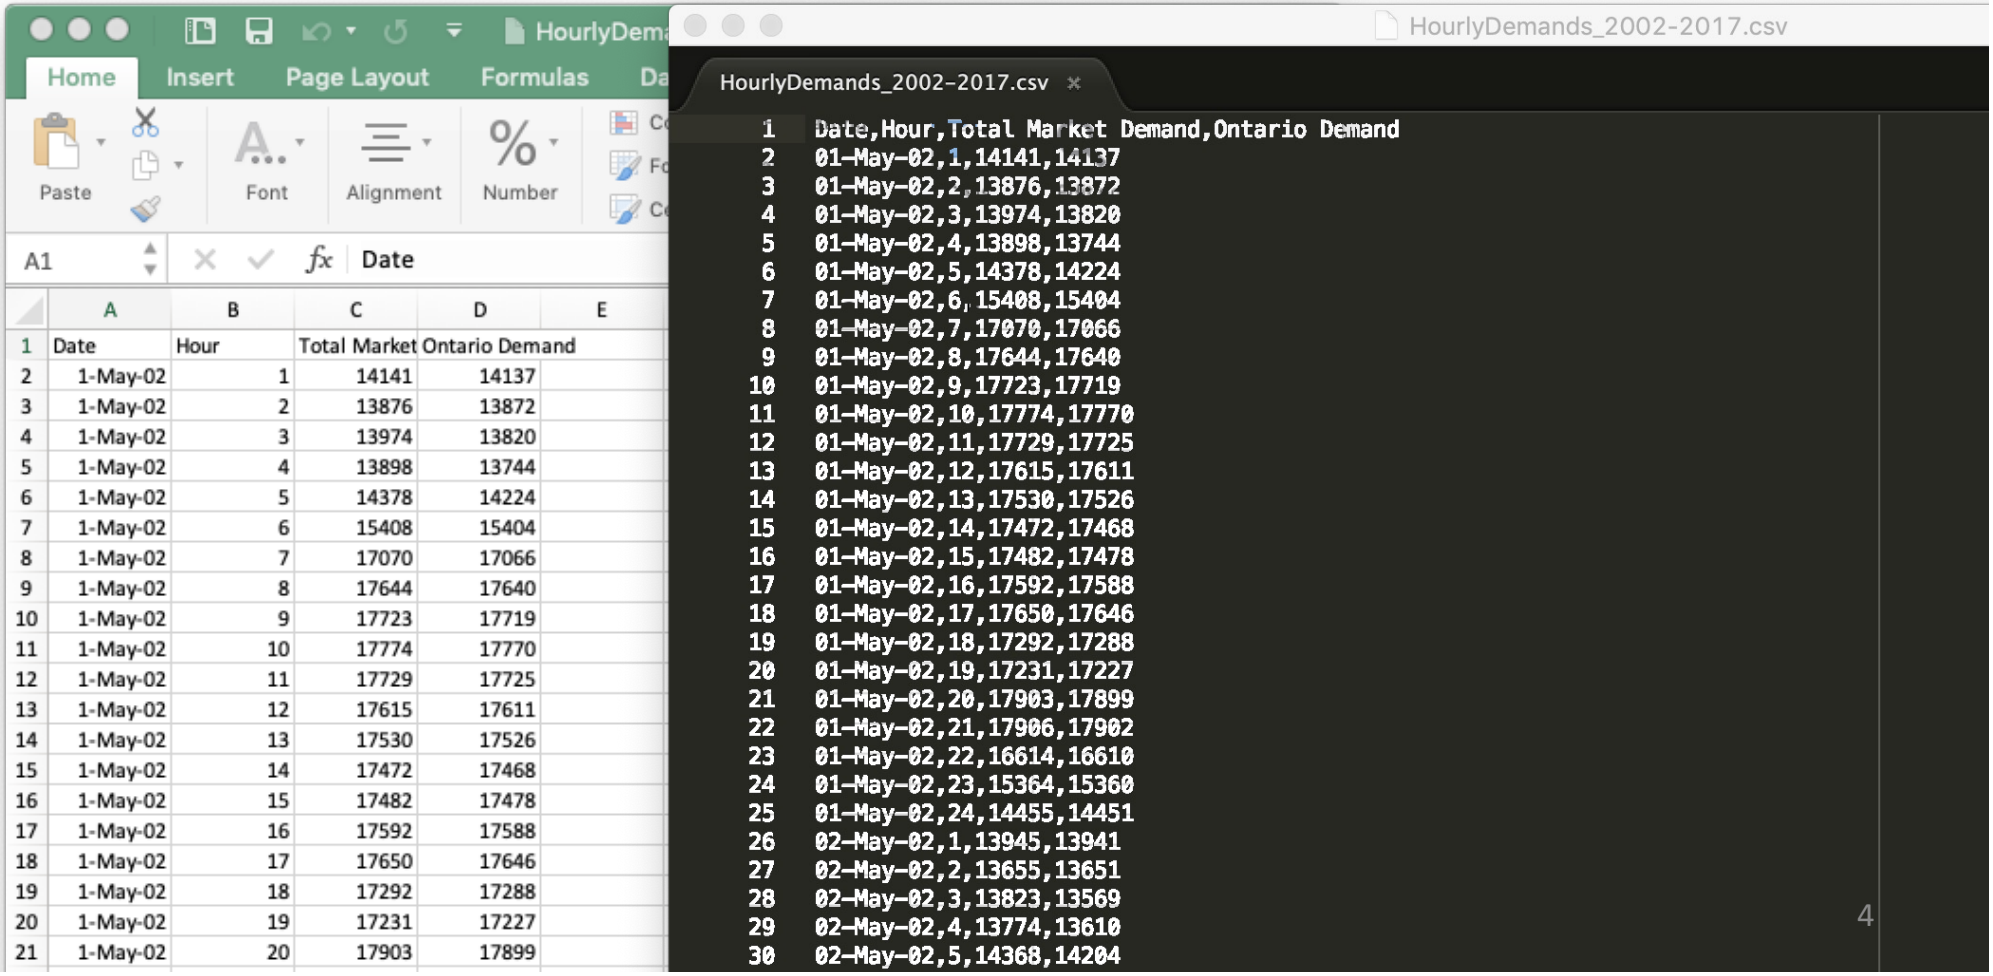

HourlyDemands_2002-2017.csv

	A	B	C	D	E
1	Date	Hour	Total Market	Ontario Demand	
2	1-May-02	1	14141	14137	
3	1-May-02	2	13876	13872	
4	1-May-02	3	13974	13820	
5	1-May-02	4	13898	13744	
6	1-May-02	5	14378	14224	
7	1-May-02	6	15408	15404	
8	1-May-02	7	17070	17066	
9	1-May-02	8	17644	17640	
10	1-May-02	9	17723	17719	
11	1-May-02	10	17774	17770	
12	1-May-02	11	17729	17725	
13	1-May-02	12	17615	17611	
14	1-May-02	13	17530	17526	
15	1-May-02	14	17472	17468	
16	1-May-02	15	17482	17478	
17	1-May-02	16	17592	17588	
18	1-May-02	17	17650	17646	
19	1-May-02	18	17292	17288	
20	1-May-02	19	17231	17227	
21	1-May-02	20	17903	17899	

```
Date,Hour,Total Market Demand,Ontario Demand
01-May-02,1,14141,14137
01-May-02,2,13876,13872
01-May-02,3,13974,13820
01-May-02,4,13898,13744
01-May-02,5,14378,14224
01-May-02,6,15408,15404
01-May-02,7,17070,17066
01-May-02,8,17644,17640
01-May-02,9,17723,17719
01-May-02,10,17774,17770
01-May-02,11,17729,17725
01-May-02,12,17615,17611
01-May-02,13,17530,17526
01-May-02,14,17472,17468
01-May-02,15,17482,17478
01-May-02,16,17592,17588
01-May-02,17,17650,17646
01-May-02,18,17292,17288
01-May-02,19,17231,17227
01-May-02,20,17903,17899
01-May-02,21,17906,17902
01-May-02,22,16614,16610
01-May-02,23,15364,15360
01-May-02,24,14455,14451
02-May-02,1,13945,13941
02-May-02,2,13655,13651
02-May-02,3,13823,13569
02-May-02,4,13774,13610
02-May-02,5,14368,14204
```

What kind of data?

- Relational data - Databases

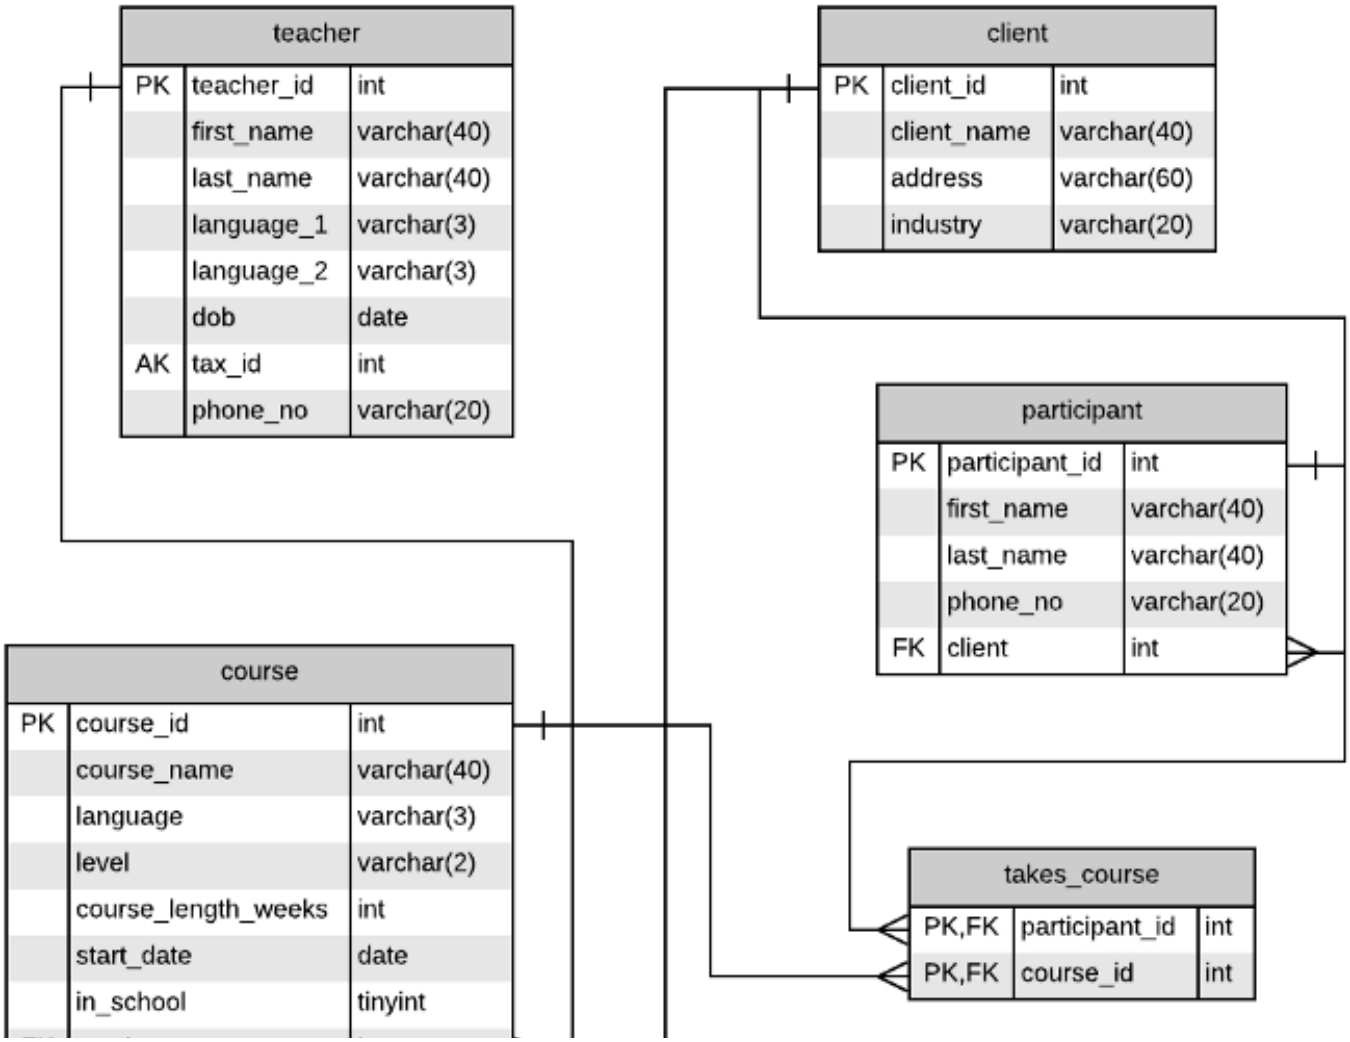

What kind of data?

- Time Series Data

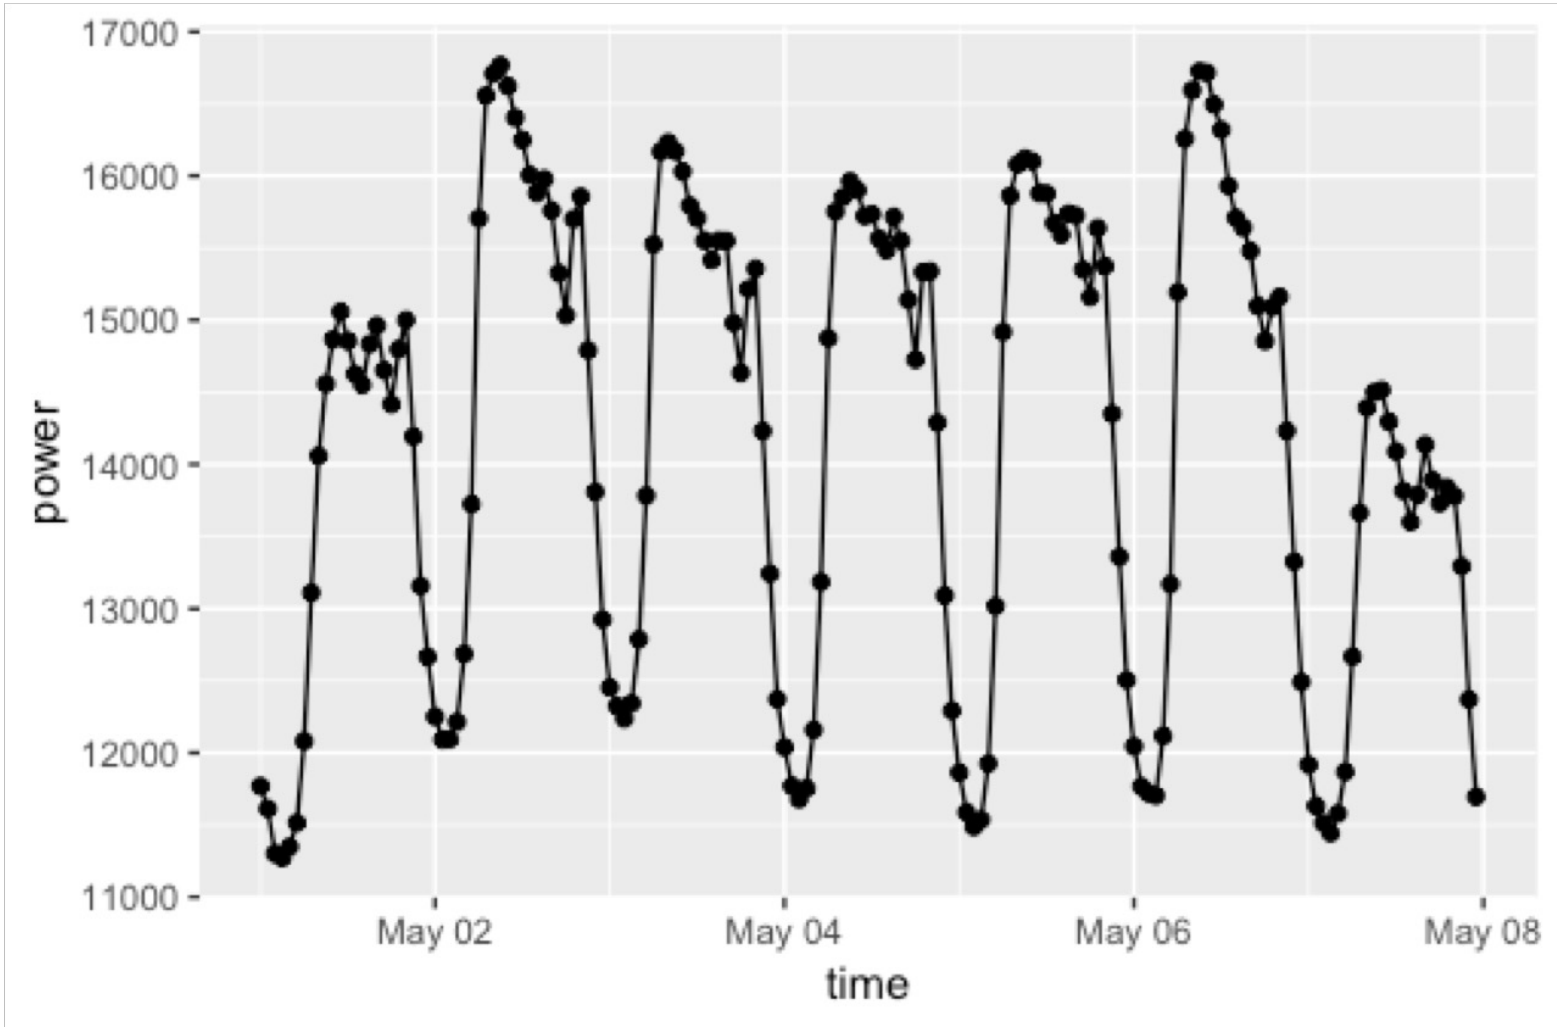

What kind of data?

- Text files, document collections

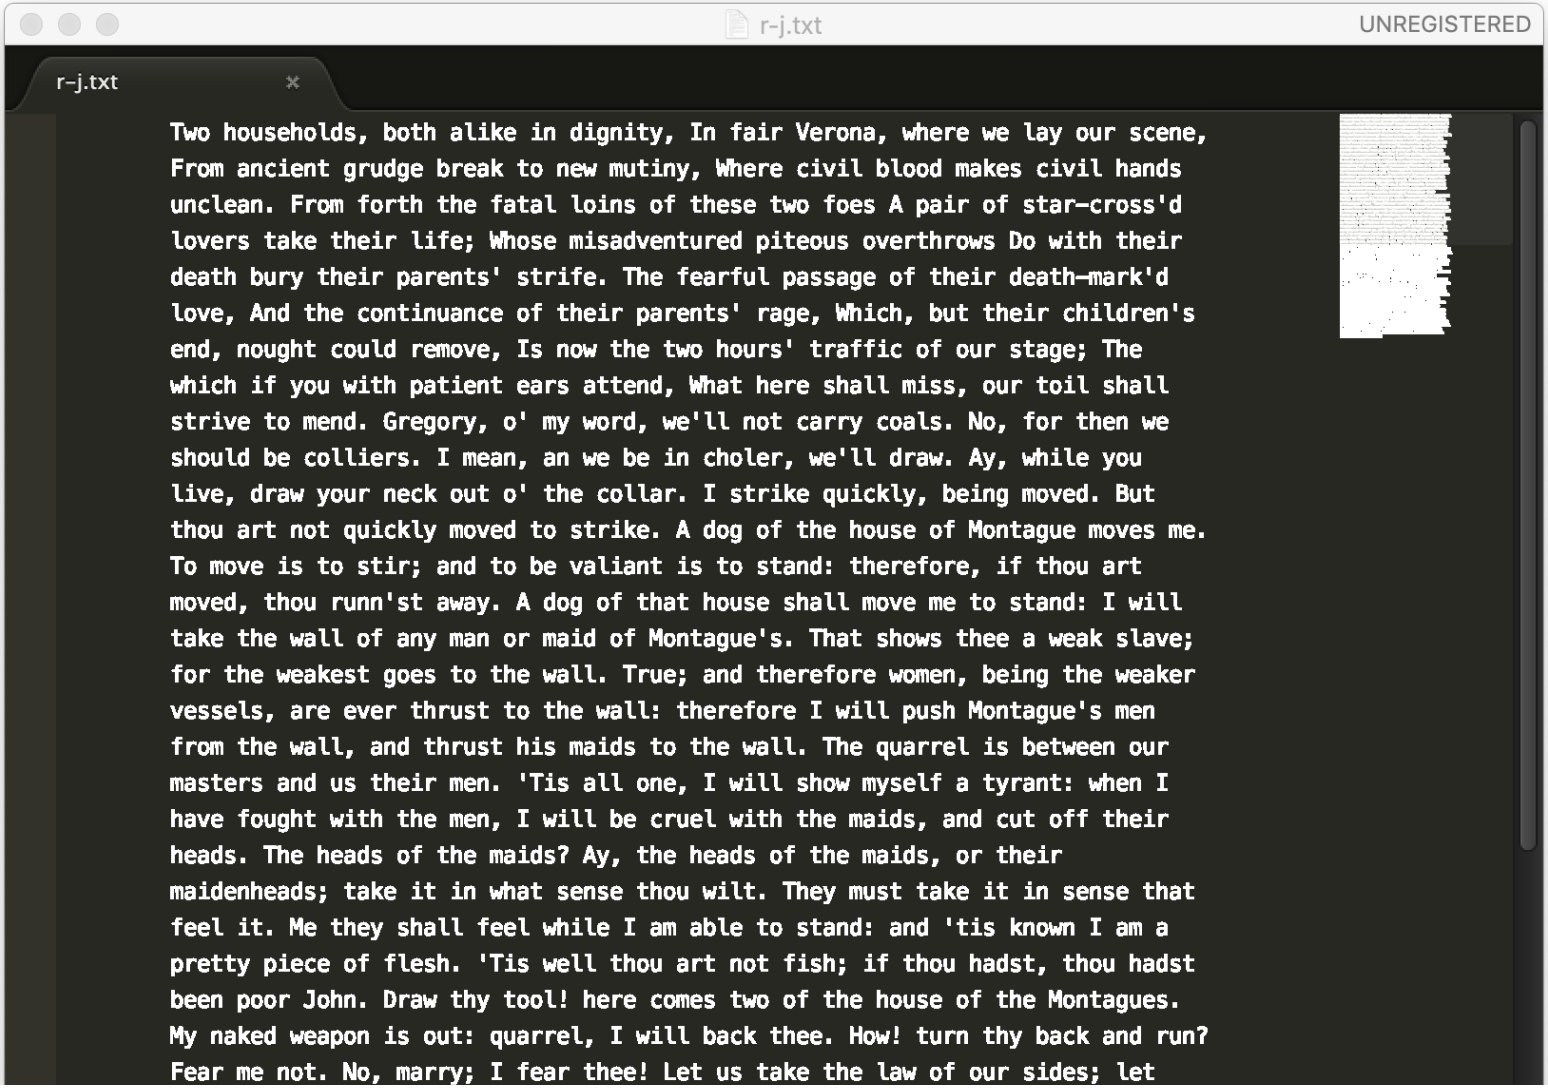

What kind of data?

- Image data

What kind of data?

- Spatial Data

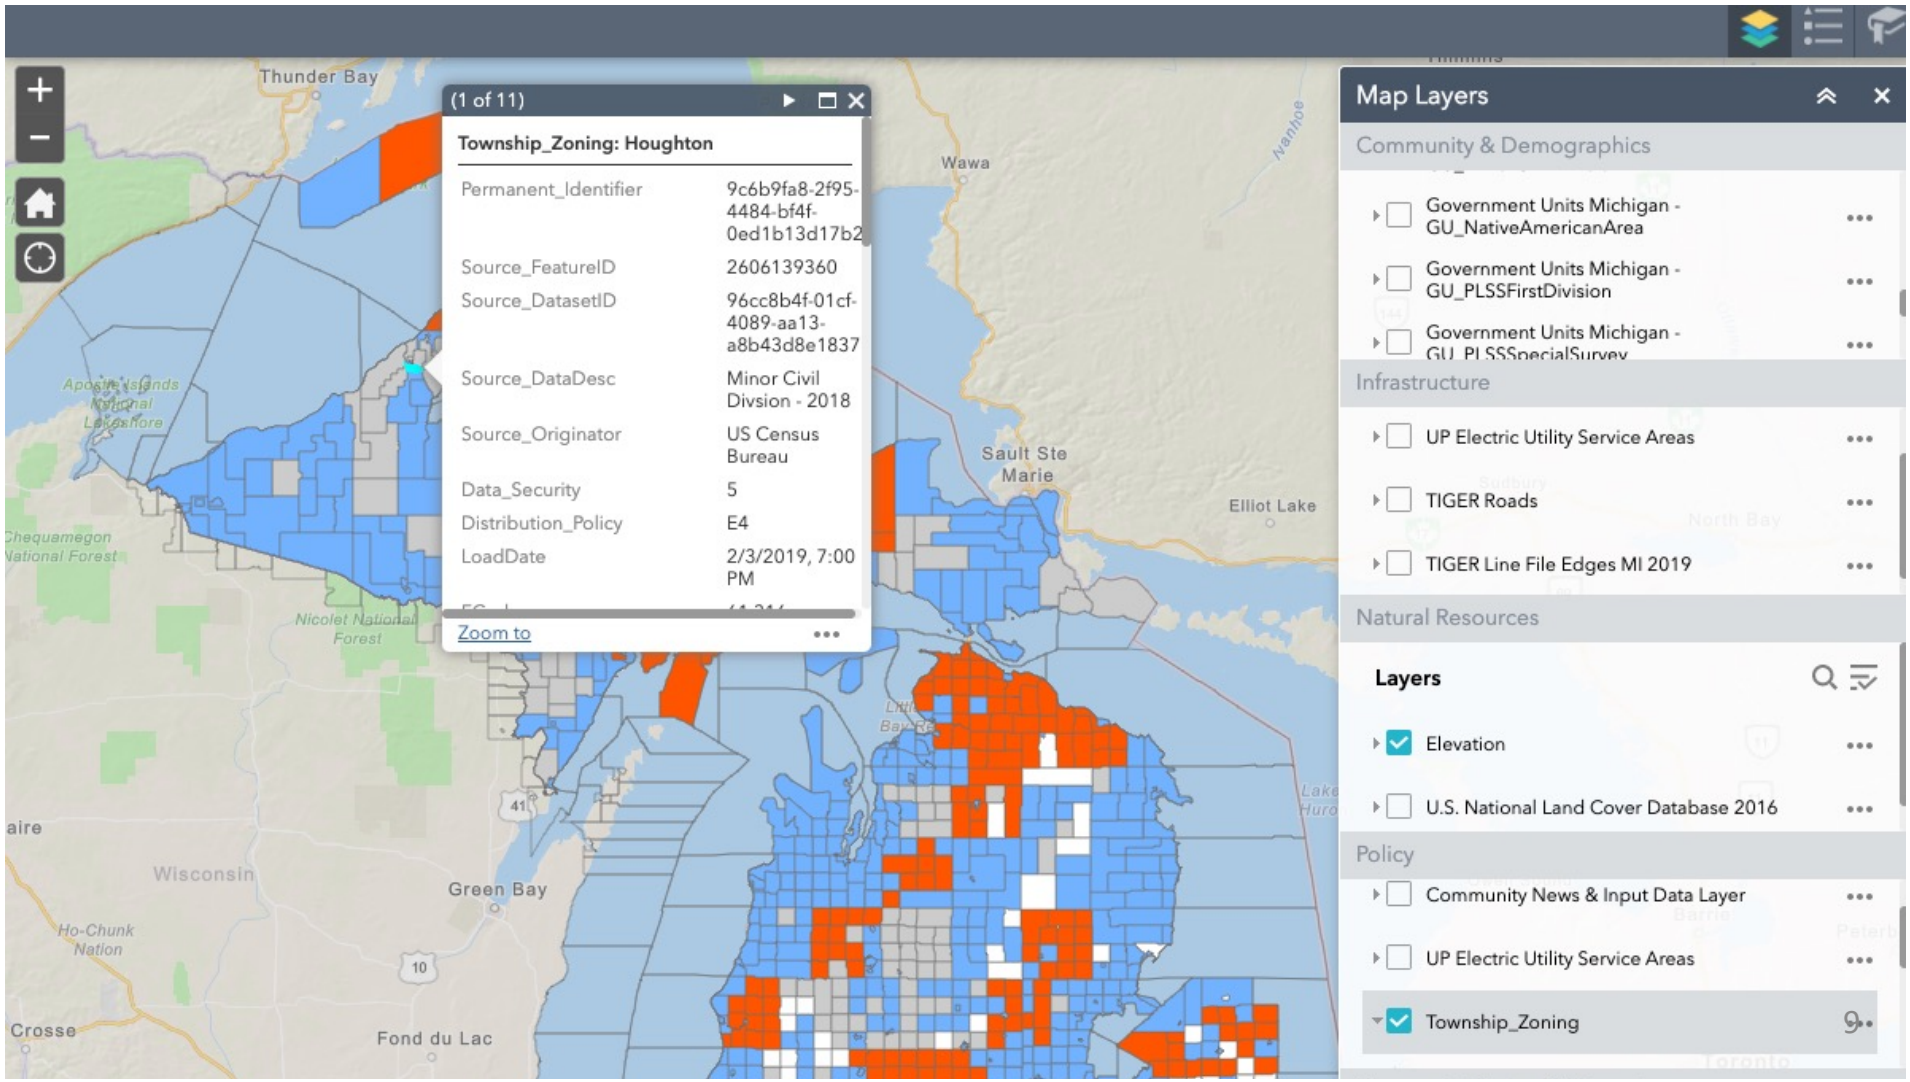

What kind of data?

- Transaction Data

TID	Items
1	Bread, Coke, Milk
2	Beer, Bread
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Coke, Diaper, Milk

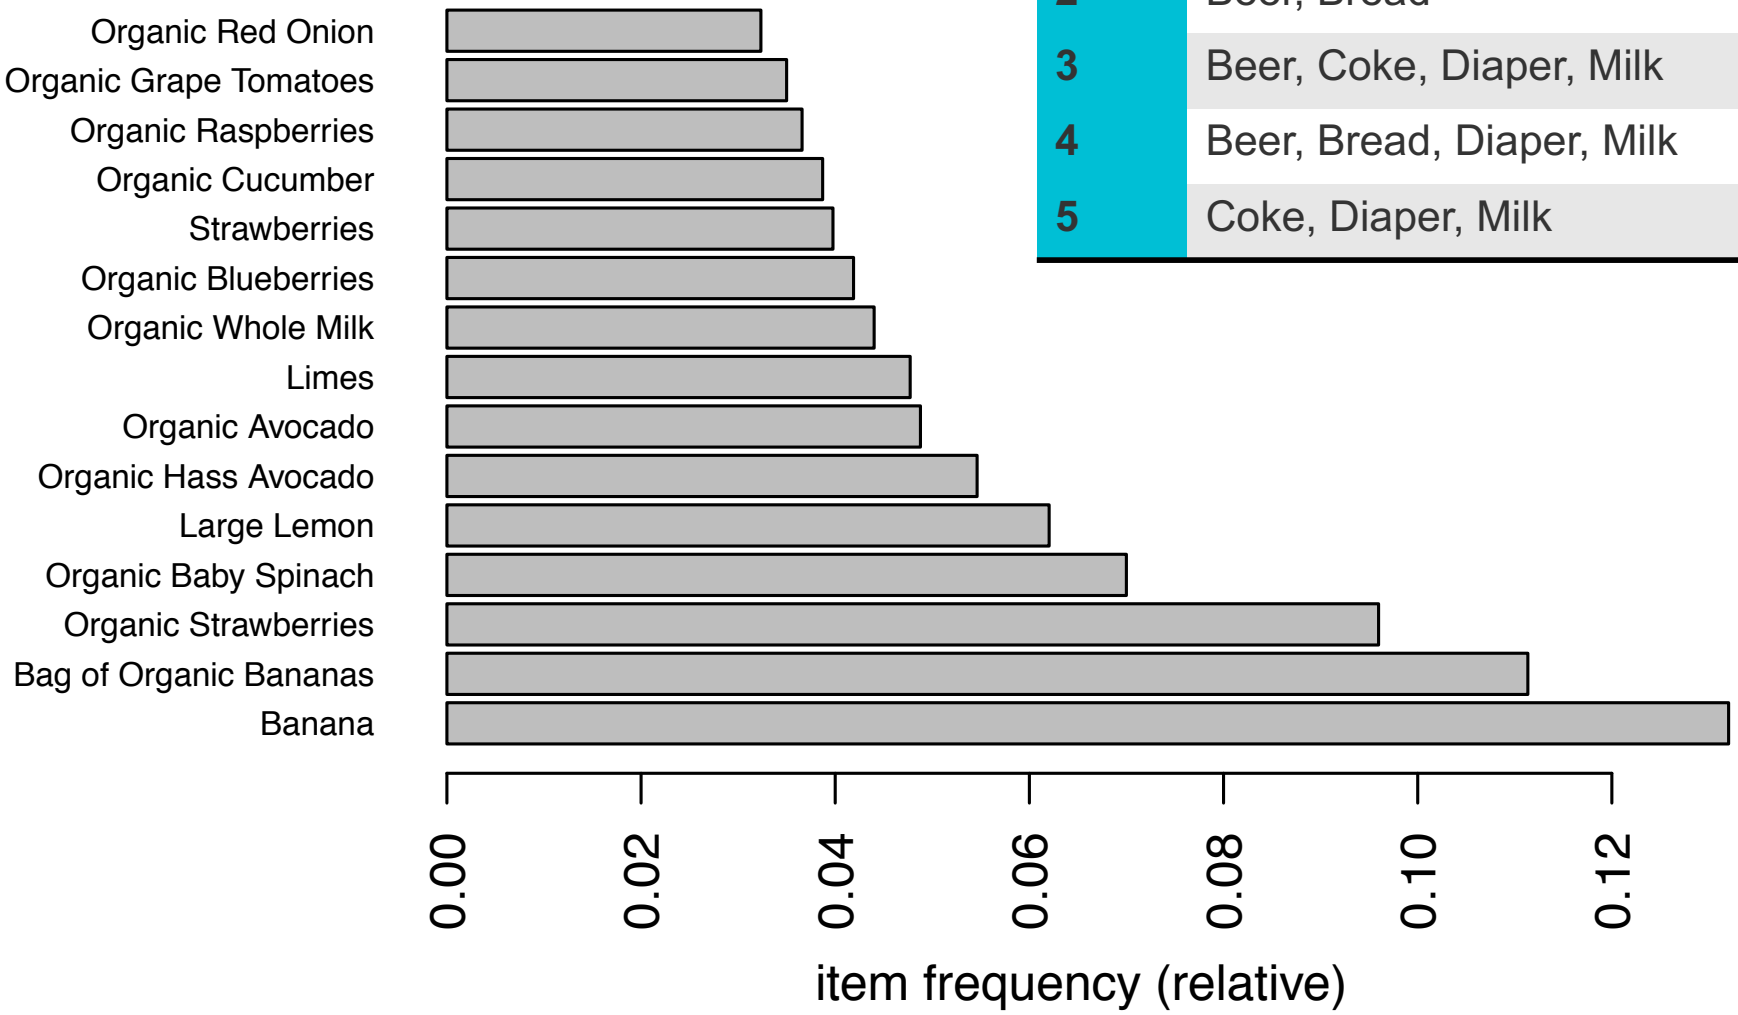

What kind of data?

- Ratings Data

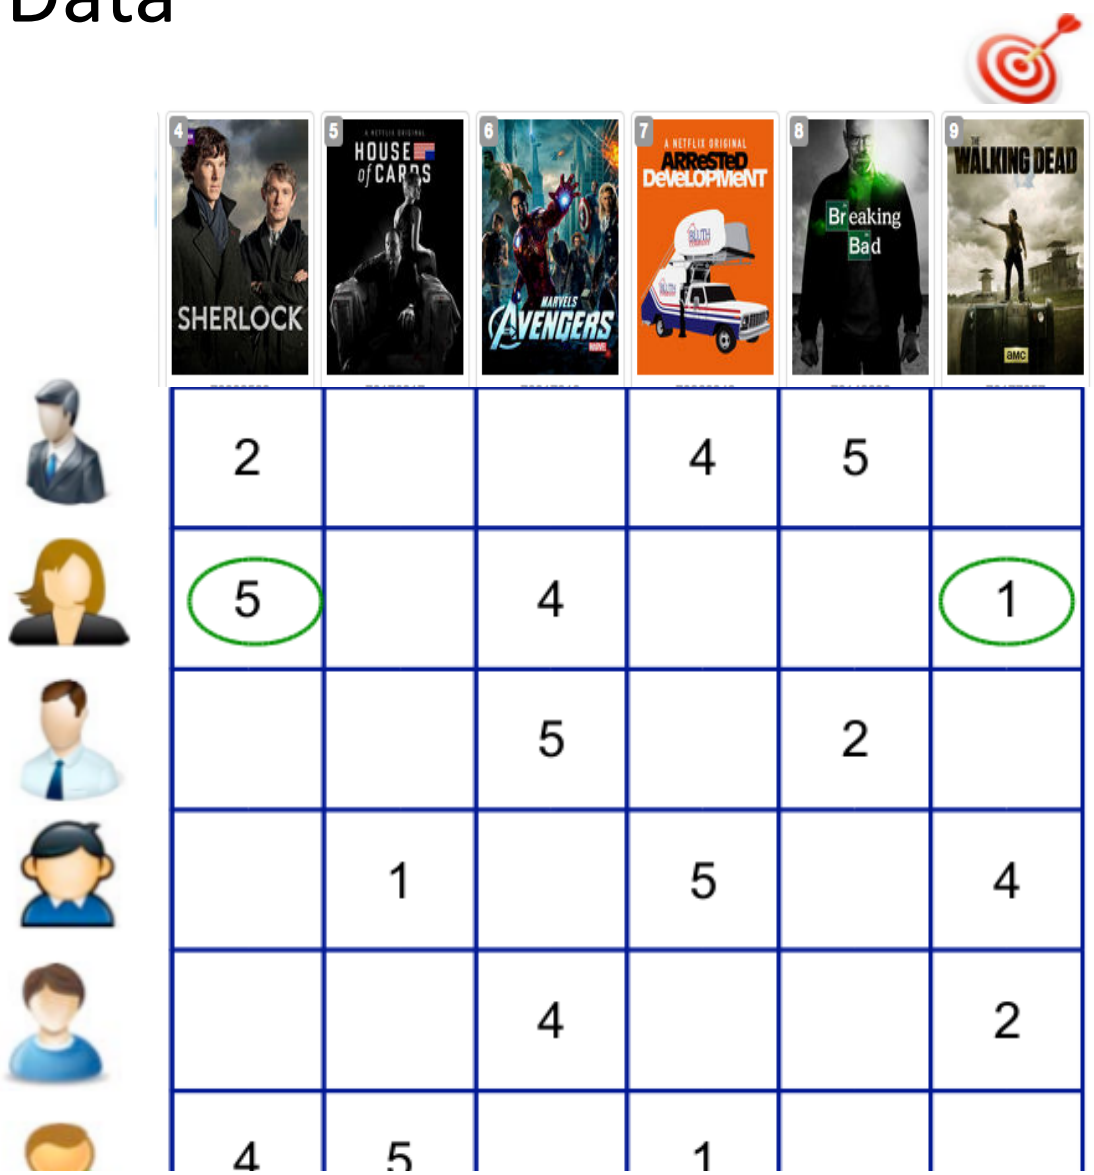

What kinds of data?

- Flat File, Vector data
- Relational data
- Text files, document collections
- Time series data
- Spatial data
- Spatio-temporal data
- Transactional data
- Ratings data
- Network data
- Image data
- Custom data for particular application
- …, and many more

What types of patterns?

- Prediction Methods
  - Use some variables to predict unknown or future values of other variables
- Descriptive Methods
  - Find human-interpretable patterns that describe the data

What type of patterns?

- Generalization / Characterization
- Classification
- Association and Correlation Analysis
- Cluster Analysis
- Recommender Systems
- Structure / Network Analysis
- Outlier Analysis
- Sequential Pattern Analysis
- …

What kinds of techniques?

- Confluence of techniques and disciplines
  - Statistics
  - Machine Learning
  - Database technology
  - Algorithms
  - Pattern Recognition
  - Parallel Computing
  - Visualization

What kinds of applications?

- Web page analysis: from web page classification, clustering to PageRank & HITS algorithms
- Collaborative analysis & recommender systems
- Basket data analysis to targeted marketing
- Biological and medical data analysis: classification, cluster analysis (microarray data analysis), biological sequence analysis, biological network analysis
- Data mining and software engineering (e.g., IEEE Computer, Aug. 2009 issue)
- From major dedicated data mining systems/tools (e.g., SAS, MS SQL-Server Analysis Manager, Oracle Data Mining Tools) to invisible data mining

Example: Fraud Detection

- Which credit card transactions are fraudulent?
- Goal: Predict fraudulent cases in credit card transactions.
  - Data: credit card transactions, information on account-holder
  - Pattern: classification
    - Labeled transaction data: fair or fraud
    - Learn a model to prediction this label
  - Technique: machine learning

Example: Fraud Detection (2)

- Which credit card transactions are fraudulent?
- Goal: Predict fraudulent cases in credit card transactions.
  - Data: credit card transactions, information on account-holder
  - Pattern: outlier analysis
    - Fraud events are rare (hopefully!)
    - Detect transaction as unusual from prior historical data
  - Technique: machine learning, statistics

Example: Clustering

- How should a set of images be placed into groups?
  - Data: image files
  - Pattern: cluster analysis
    - Group images by similarity
    - How to measure similarity?
  - Techniques: machine learning

In this class …

- You will study algorithms that exploit and reveal patterns in data.
- Goal:
  - You will learn how to think about problems in data mining
  - You will learn about a set of data analysis tools:
    - How to use them
    - What their assumptions are
    - The capabilities and limitations

Methods to Be Examined

- Supervised learning – Classification
- Unsupervised learning – Data Reduction
- Text Mining
- Unsupervised learning – Clustering
- Association Mining
- Recommendation Systems
- Web Mining (if time permits)

Popular Tools for Data Mining

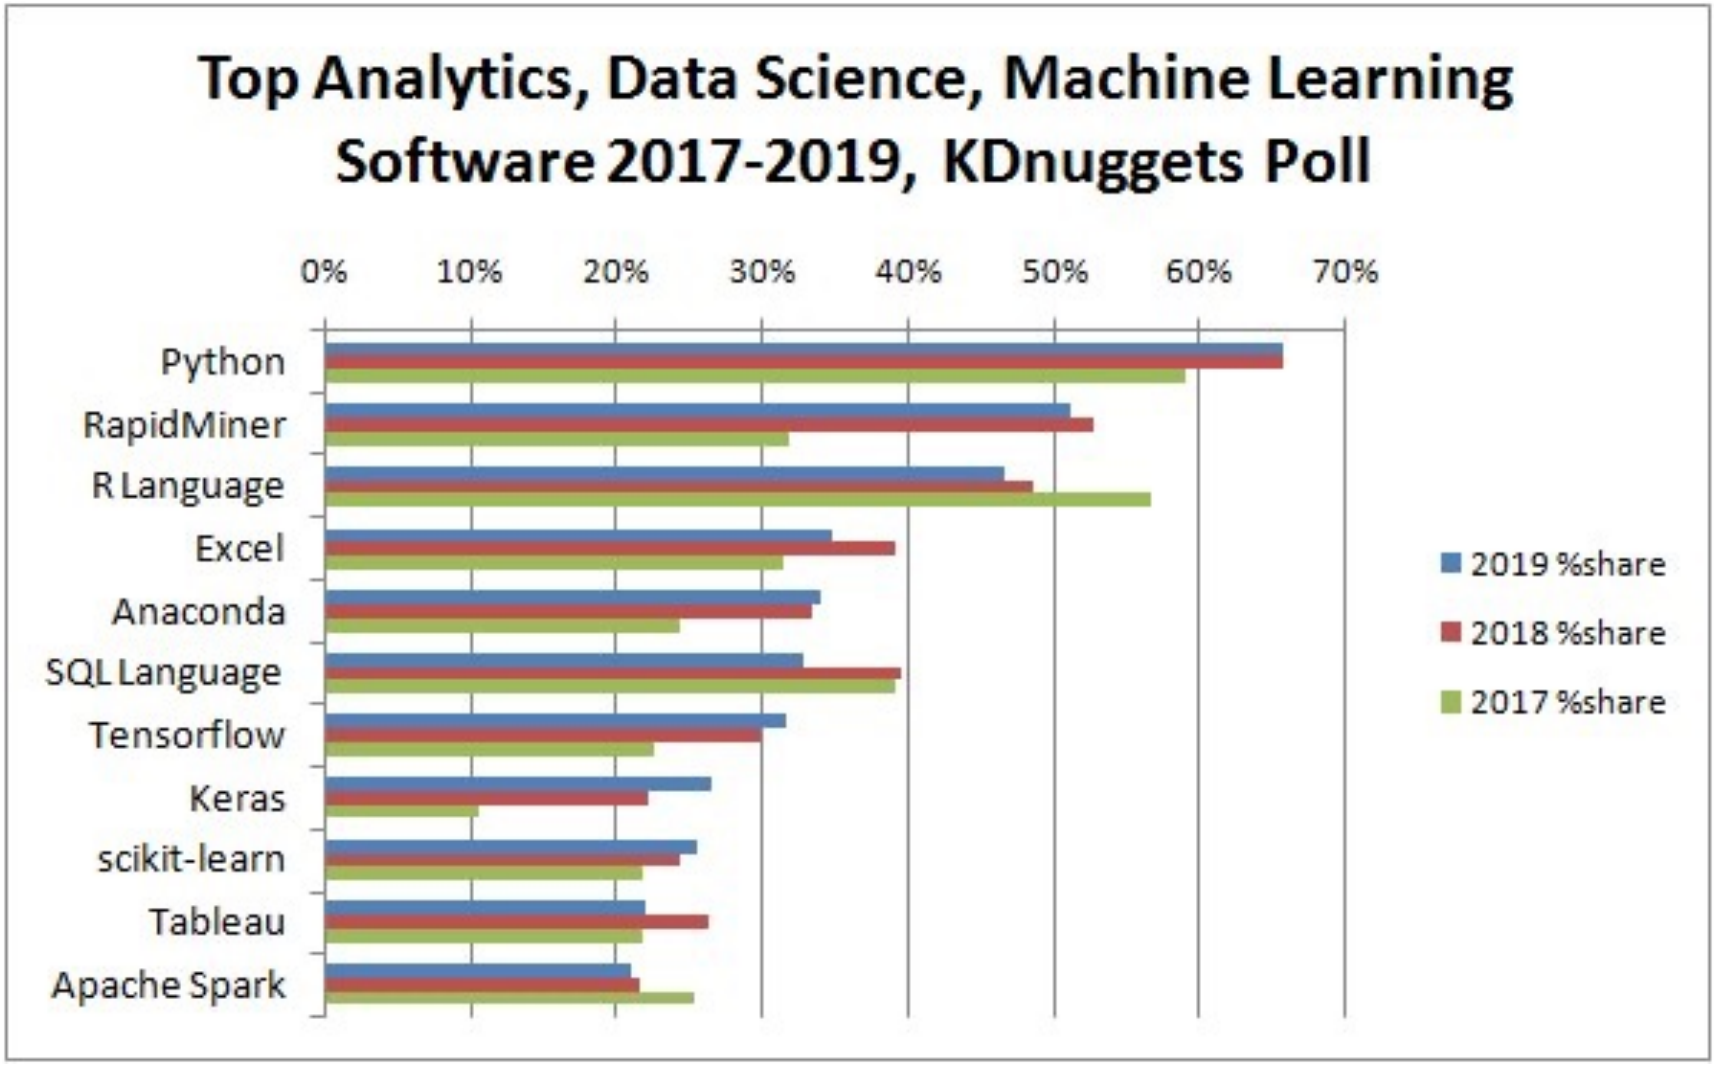

https://www.kdnuggets.com/polls/

Types of Tools

Simple graphical user interface

Process oriented

Programming oriented

Tools: Simple GUI

- Weka: Waikato Environment for Knowledge Analysis (Java API)
- Rattle: GUI for Data Mining using R

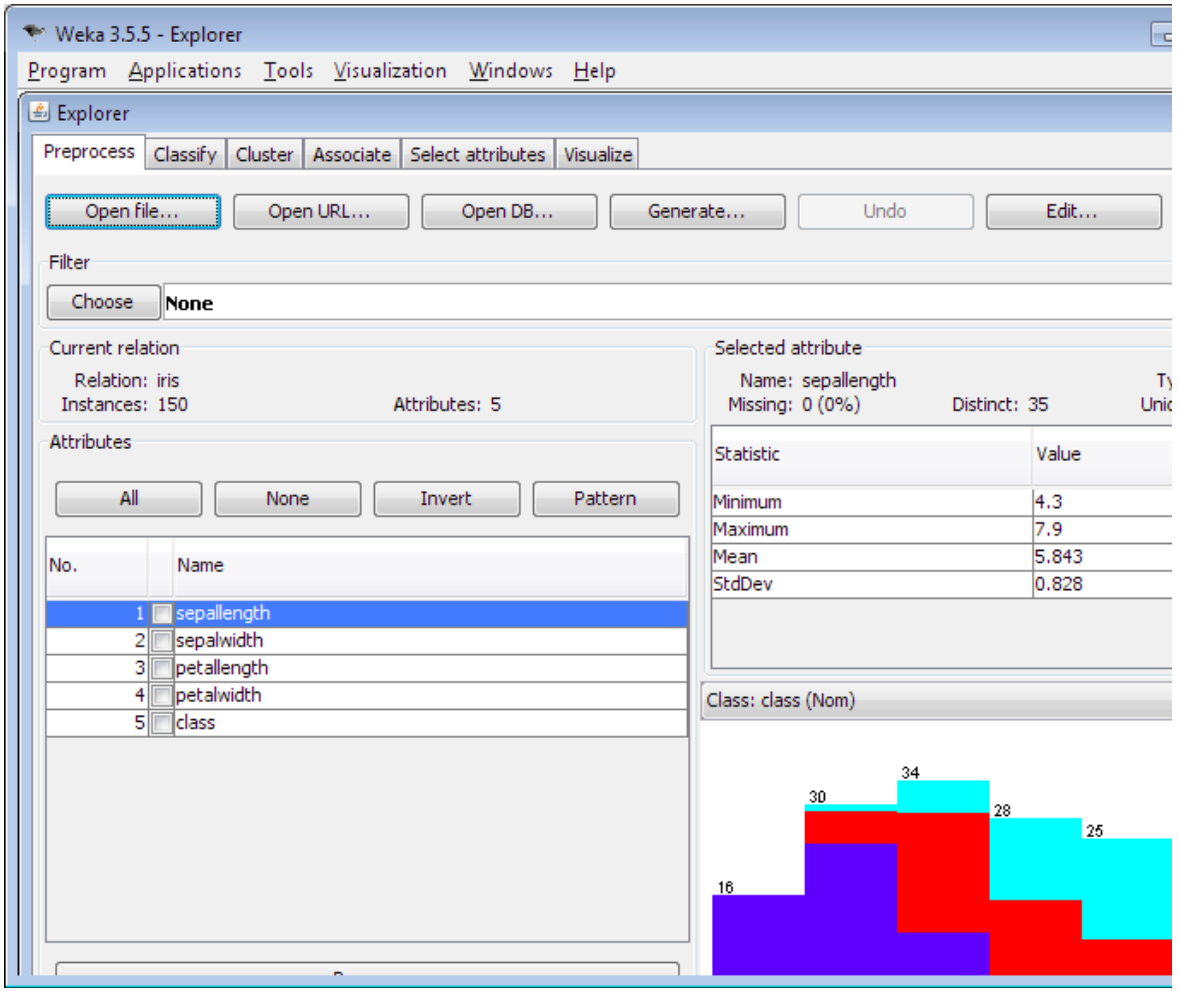

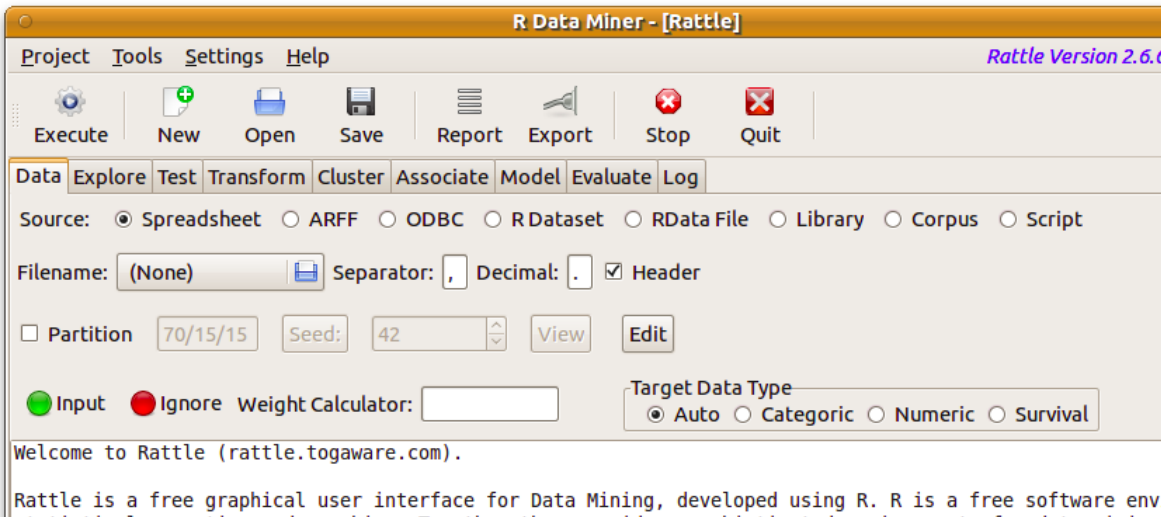

Tools: Process oriented

- SAS Enterprise Miner
- IBM SPSS Modeler
- RapidMiner
- Knime
- Orange

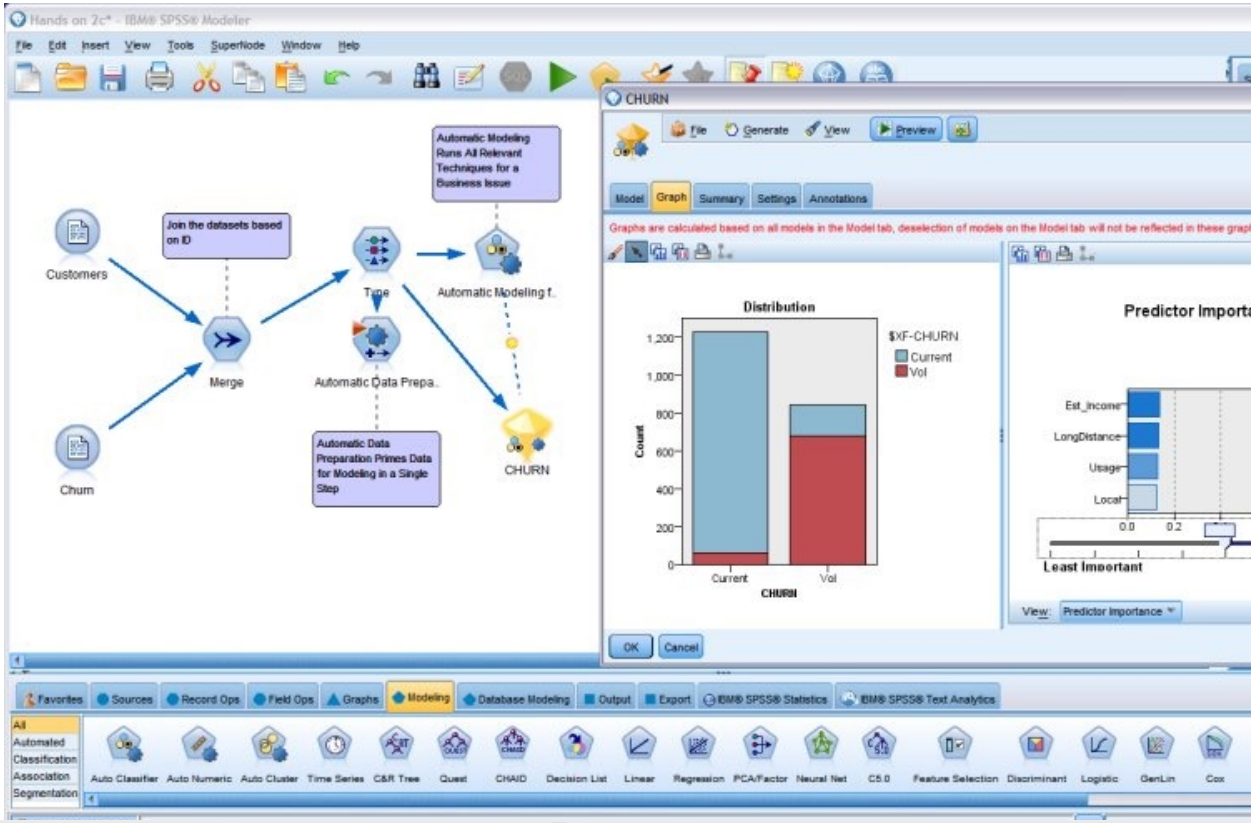

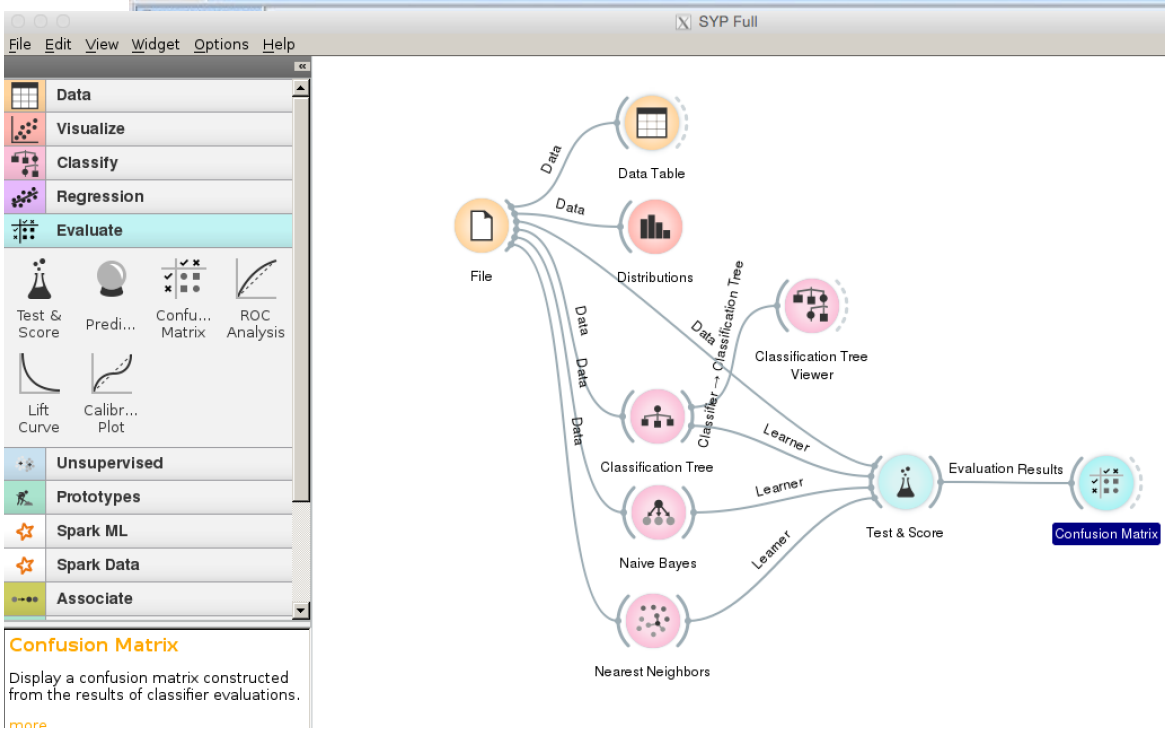

Tools: Programming oriented

- R
  - RStudio IDE
    - Caret
    - Tidyverse
- Python
  - Jupyter notebooks
    - Numpy
    - Pandas
    - Scikit-learn

Data Mining vs. Privacy

- Tension between data mining and personal privacy

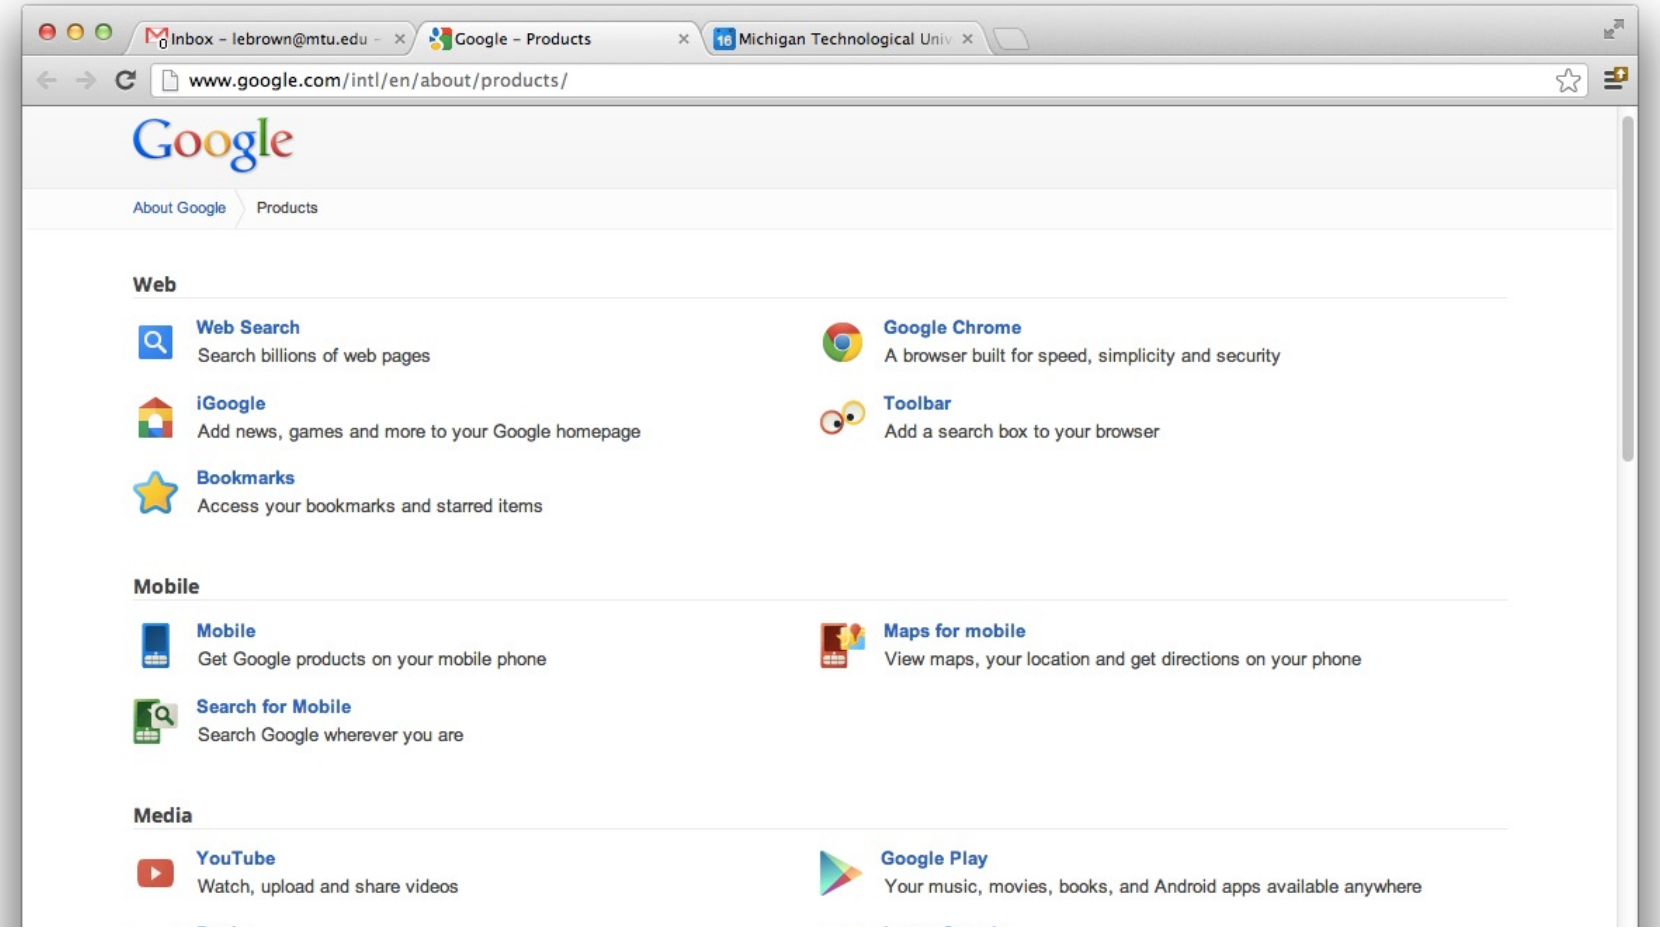

Data Mining vs. Privacy

- How can we leverage sensitive personal data for research / commercial purposes?

- 3 cases
  - AOL search data set
  - Netflix prize
  - Barabasi mobile study

Examples from C. Volinsky

Case 1: AOL Search Data

- Aug. 4, 2006 – AOL releases 20M search terms by anonymized users “for research purposes”
- Within hours, uproar on blogs
  - “The utter stupidity of this is staggering” – TechCrunch
- Aug. 7, 2006 – AOL removes data, issues apology
  - “this was a screw-up, and we are angry”
  - “an innocent enough attempt to reach out the the research community”
- Aug. 9, 2006 – NYT front page story
  - Identifies user
- Aug. 21, 2006 – CTO resigns

Case 1: AOL Search Data

- What's the big deal?
  - How and why people search is often personal and may contain information they do not want released to public
- What went wrong?
  - Not well thought out
  - Poor internal controls
  - Lack of understanding on anonymizing
- Fallout
  - CTO + at least two others fired
  - Data is still out there

Case 2: Netflix Prize

- Oct. 2006: Netflix released anonymized movie ratings from its customer database
  - 100M ratings, 500K customers (<10% of data)
  - Random integer for user ID
  - "some of the rating data for some customers in the training and qualifying sets have been deliberately perturbed in one or more of the following ways: deleting ratings; inserting alternative ratings and dates; and modifying rating dates"
- 2007, Paper claiming to de-anonymize Netflix data

Case 2: Netflix Prize

- Narayanan and Shmatikov
  - "The adversary with a small amount of background knowledge about an individual … can identify with high probability that individual's record in the data and learn … sensitive attributes"
  - Claim Netflix's data sanitization not relevant
  - Basic Idea:
    - With aux info on 8 movies, where 2 can be wrong, and dates are known within 14 days, 99% de-anonymization
    - Aux info can come from other web-sites (IMDB), personal contact, etc.

Case 2: Netflix Prize

- Much ado about nothing
  - Paper is technically correct, but dates are key
  - Without dates, you must know 8 movies, all outside the top 500 to get over 80% chance of de-anonymization
  - Aux info is not easy to come by for many people
  - No identities released
- Netflix did it right
  - Consulted with top machine learning experts
  - Gained new knowledge in machine learning and also privacy fields
- Fallout
  - Netflix was planning another challenge was canceled to due privacy concerns

Case 3: Barabasi Mobile Study

Gonzalez, Hidalgo, and Barabasi (2008)

- Article in Nature outlines study on human mobility patterns
  - 100000 individuals selected randomly from dataset of 6 million
  - Unidentified country (unclear if researchers knew)
  - Cell tower location at start of call
  - 206 individuals were “pinged” every two hours for a week
- Findings
  - “humans follow simple, reproducible patterns”
  - Sample finding: Nearly three-quarters of those studied mainly stayed in a 20-mile circle for half a year.
  - Results “could impact all phenomena driven by human mobility, from epidemic prevention to emergency response and urban planning”

Case 3: Barabasi Mobile Study

- Uproar over “secret tracking” of cell phone users
  - Blowback of negative feedback to Nature and scientists
  - Study would be “illegal in the US”
  - Approval from ONR review board and Northeastern review board. Barabasi did not check with an “ethics panel”
- Response
  - Hidalgo: “the data could be misused”, but we were “not trying to do evil things. We are trying to make the world a little better.”
  - Northeastern and Nature backed the research
  - Continues to be referenced as an example of dangerous research
  - Risk and reward both very high

Data Mining and Ethics

- Privacy is not the only issue in data mining
  - Selling of Data
  - Transparency of Data Collection
  - Security
  - Discrimination and Bias

- We will explore these topics periodically throughout the semester

Data Mining Introduction Summary

Data Mining is interdisciplinary and overlaps significantly with many fields

- Statistics
- CS (machine learning, AI, data bases)
- Optimization

Data Mining requires a team effort with members who have expertise in several areas

- Data management
- Statistics
- Programming
- Communication
- + Application domain