

Data Mining: Classification Ensemble Methods

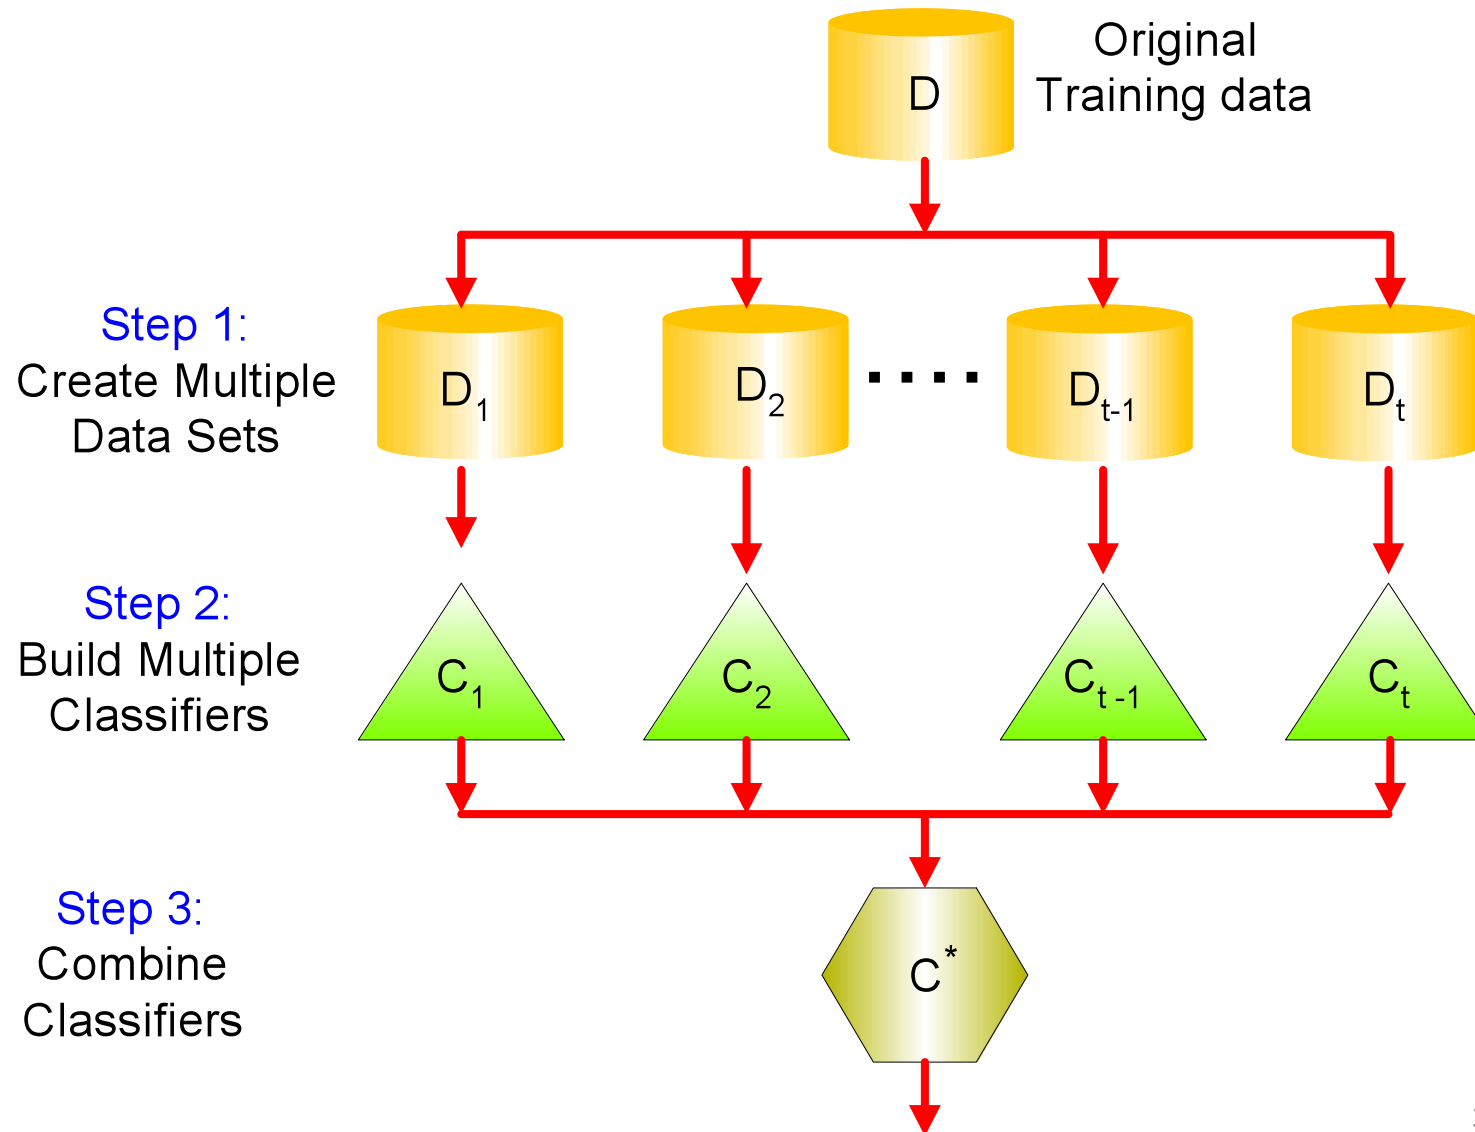
Laura Brown

Some slides adapted from P. Smyth; A. Moore; D. Klein; R. Tibshirani
Han, Kamber, & Pei; Tan, Steinbach, & Kumar; L. Hannah;
J. Taylor and L. Kaebling

Basic Ensembles

- Build different models on data
 - E.g., build same model multiple times with different random seeds or different hyperparameters
- Average/Majority of model results
 - Soft voting: averaging the predicted probabilities and take the arg max
 - Hard voting: use each model's prediction and select most commonly predicted

Ensemble Methods



Ensemble Learning

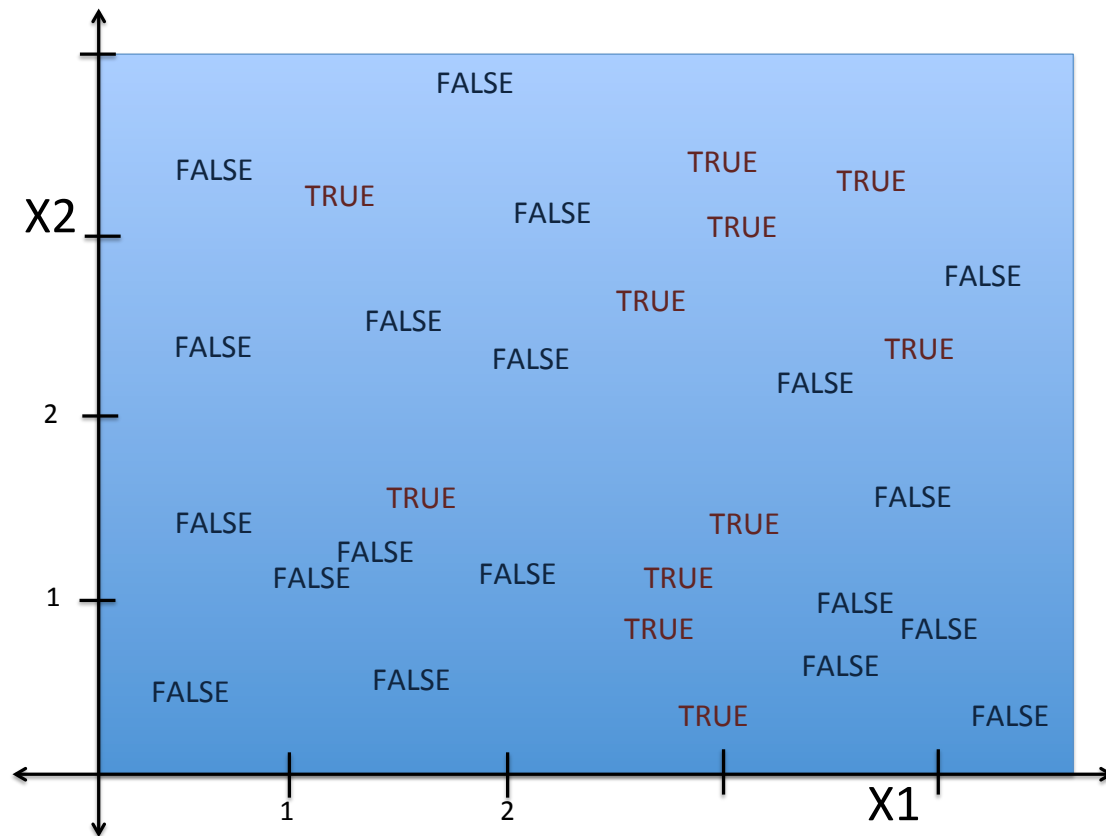
- Select a multitude of hypotheses to make predictions and combine their results
 - Use a number of different learners
 - Use the same learner with different hyperparameters
- If the errors are independent (or approaching independence), then the different hypotheses are complementary
- Combinations are most likely to be right than any individual hypothesis

Why Ensemble Methods Work?

- The simple models that go into the ensemble are easy to learn
 - But, they have a limited hypothesis space
- By combining (averaging) many different simple models, the models can fit the data well and have a large hypothesis space

Example: Ensemble Method

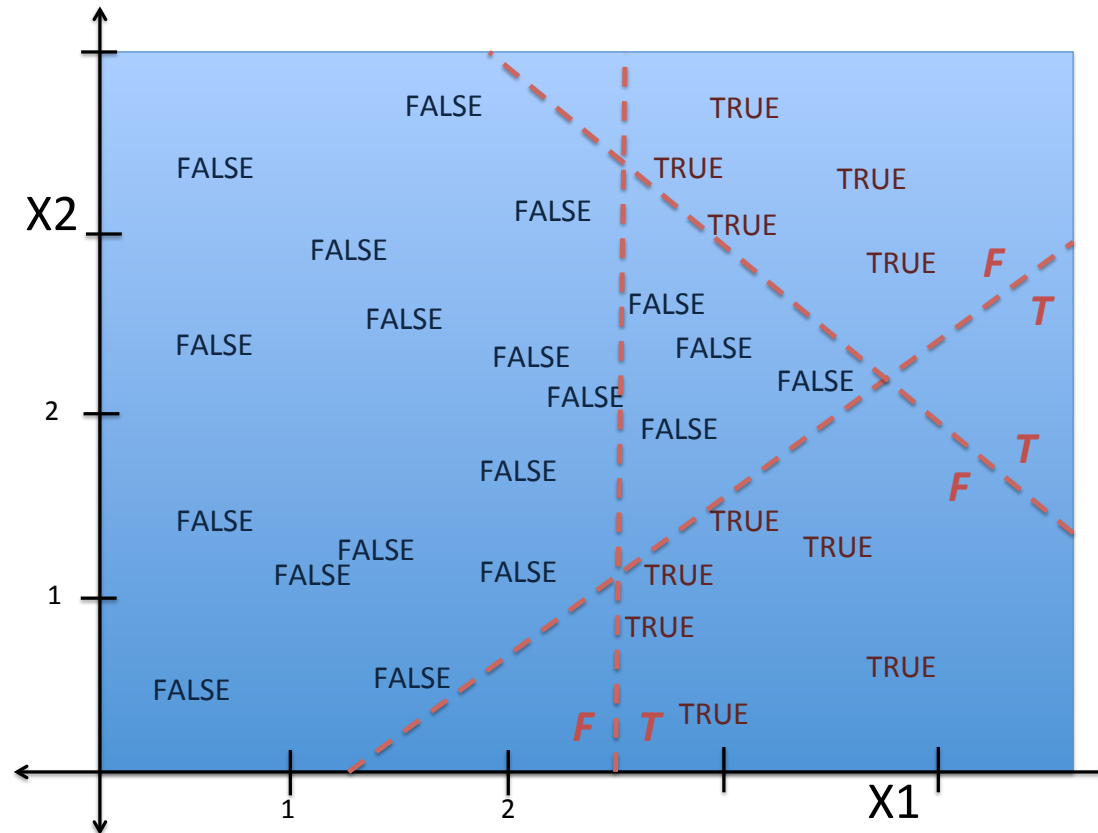
- Consider Linear Classifiers: divide space with a hyperplane



- No single hyperplane perfectly separate the two classes
- How can they be combined?

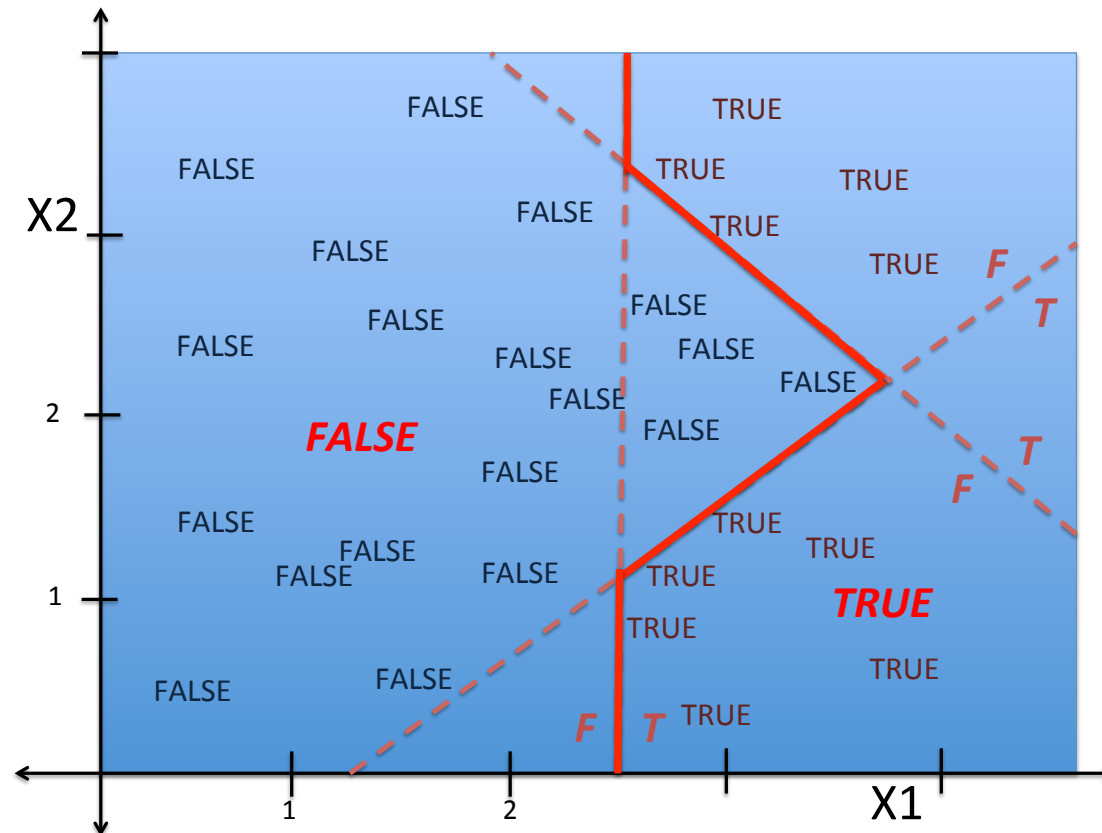
Example: Ensemble Method

- Consider Linear Classifiers: divide space with a hyperplane



- No single hyperplane perfectly separate the two classes
- How can they be combined?

Example: Ensemble Method



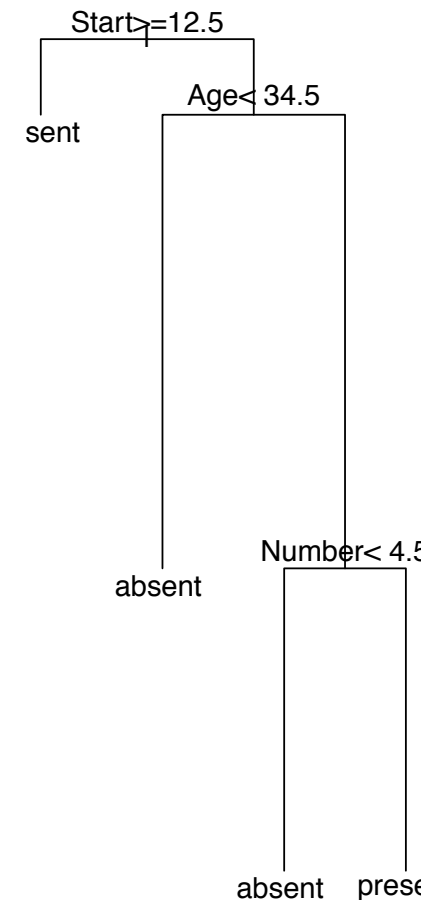
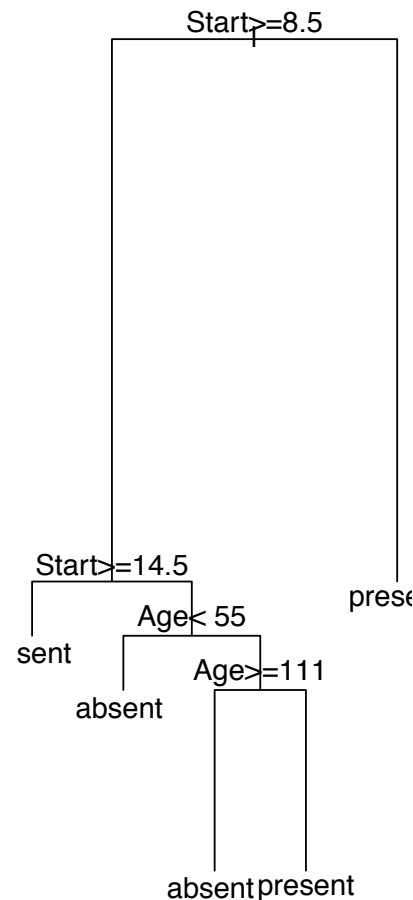
- Combine by assigning majority labels
- Results in non-linear surface

Ensemble Methods and Trees

Kyphosis data:

Trees

- Flexible models – work for both regression and classification
- Tend to fit pretty well, but do not always have best predictive error
- Trees are unstable



Ensemble Methods and Trees

- Instability
 - Small changes in data (or fitting method) produce big changes in outcome
 - This is good for ensembles!
Diverse results

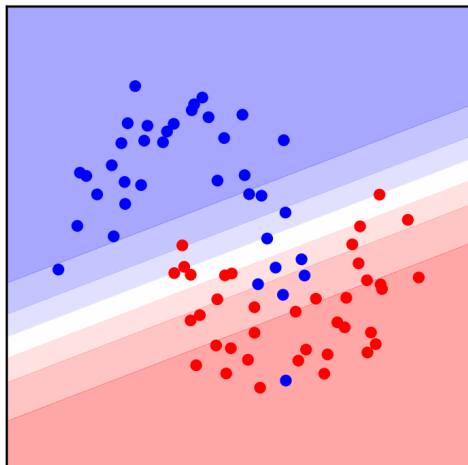
Types of Ensemble Learning

- Model Averaging flavors
 - Fully Bayesian: average over uncertainty in parameters and models
 - “empirical Bayesian”: learn weights over multiple models
 - e.g., stacking and bagging
 - Build multiple models in a systematic way and combine them
 - VotingClassifier
 - Bagging
 - Random Forests
 - Stacking / Ensemble
 - Boosting

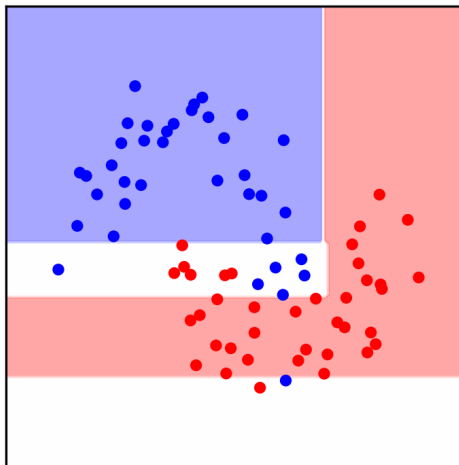
Voting Classifier

- Simplest Approach
- Learn multiple models or models with different hyperparameters on same data set
- Often better to use models that are different from each other

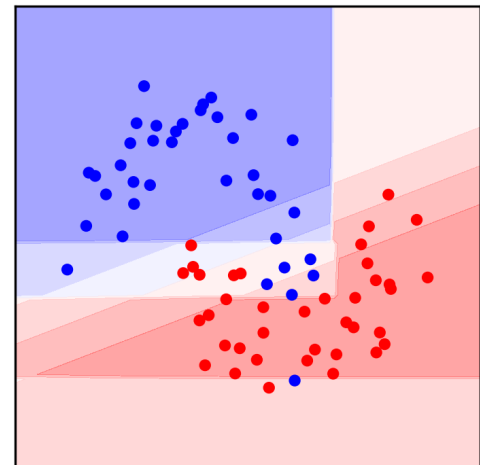
Logistic Regression acc=0.84



Decision Tree acc=0.80



Voting Classifier acc=0.88

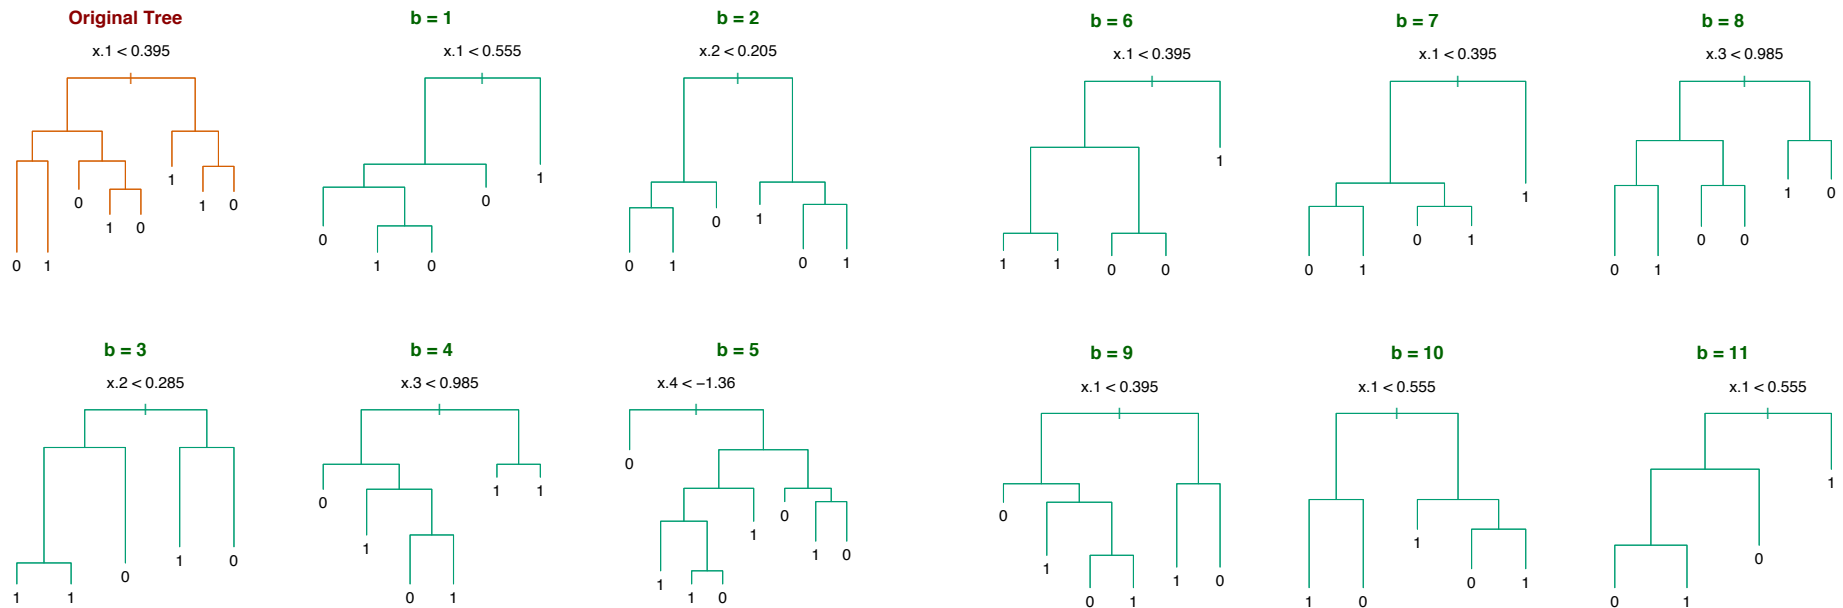


Bagging: Bootstrap Aggregation

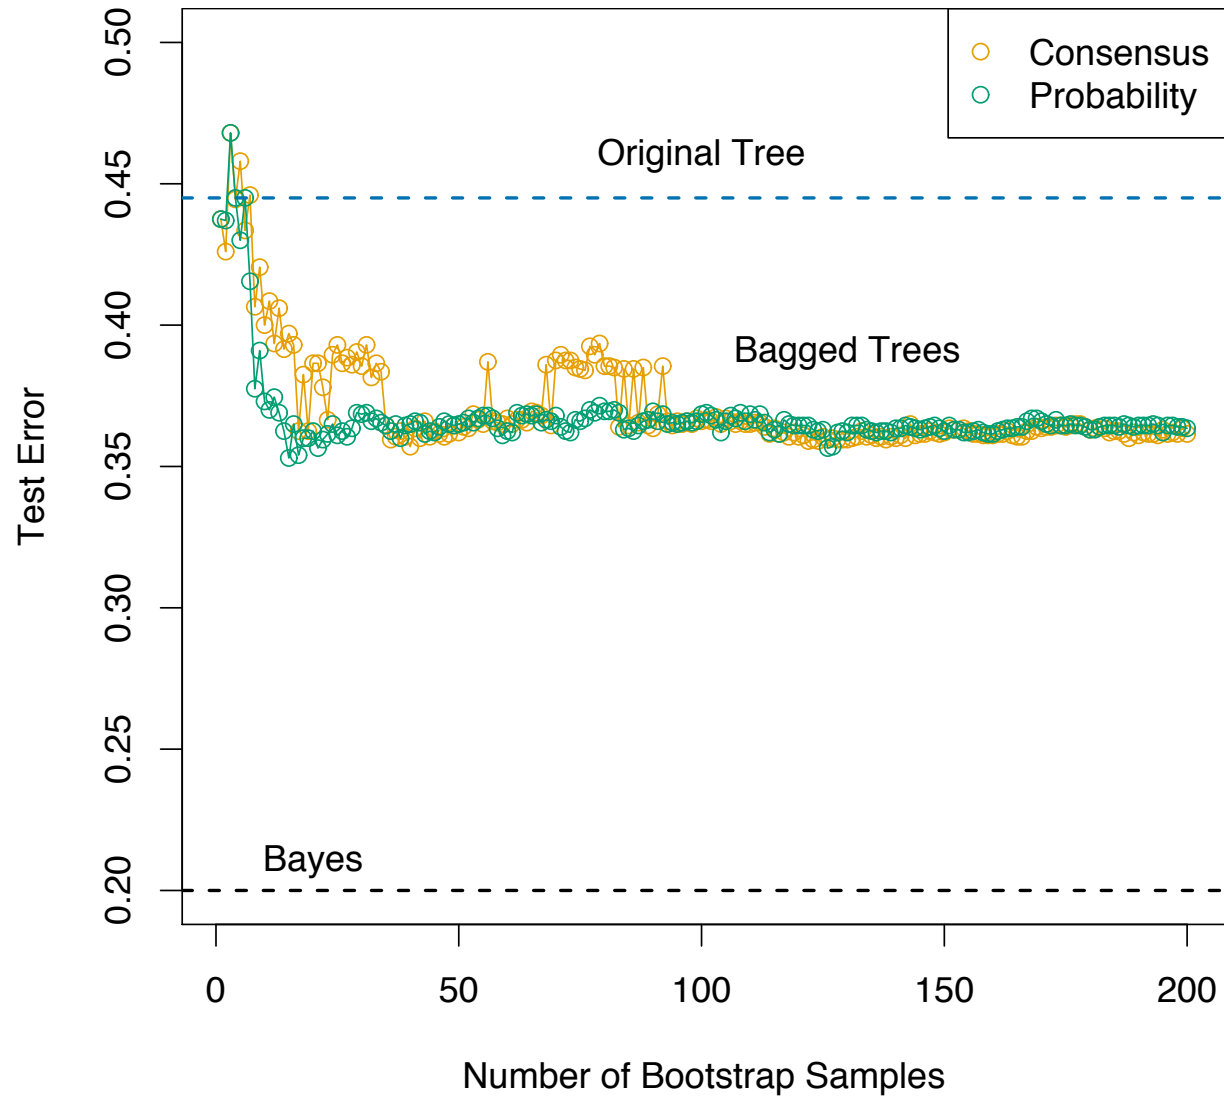
- Training:
 - given a data set D with n tuples, at each iteration i , a training data set D_i of n tuples is sampled with replacement from D (bootstrap)
 - A classifier model M_i is learned for each D_i
- Classification: classify unknown sample x
 - Each classifier M_i returns its prediction
 - The bagged classifier M^* counts/averages the votes and assigns the class to x
- Performance
 - Often better than single classifier on data D , but loses interpretability of model

Example: Simulated Data

- From ESL (8.7.1), $n=30$ training data points, $p=5$ features, and 2 classes.



Example: Simulated Data



Example: Breiman's Bagging

- From Breiman's paper: compare misclassification error of tree and that of bagging result

Data Set	\bar{e}_S	\bar{e}_B	Decrease
waveform	29.1	19.3	34%
heart	4.9	2.8	43%
breast cancer	5.9	3.7	37%
ionosphere	11.2	7.9	29%
diabetes	25.3	23.9	6%
glass	30.4	23.6	22%
soybean	8.6	6.8	21%

Bagging

- Pros:
 - Easy to implement
 - Works better than model on their own
 - Very fast and parallelizable
- Cons:
 - May not work as well as Boosting
 - Works best with high variance, low bias, low correlation estimators

Bagging

- Ensembles tend to work best with “not too complicated” hypotheses
- Why?
 - Simpler models are often less correlated
 - Cover a larger part of the hypothesis space
- Lets work again with trees, but try to de-correlate them

Random Forests

- Why de-correlate trees?
 - If inputs are the same, tree generation will produce the same branching path
 - But small changes in inputs, can lead to large changes in output
 - Force trees to split on different attributes
- **Randomly select a subset of attributes** that it can split on

Random Forests

- Grow each tree on independent bootstrap sample
- At each node:
 - Randomly select m variables out of all p (independently for each node)
 - Find the best split of selected m variables
- Grow each tree to maximum depth
- Vote / average the trees to get predictions

Hyperparameters for RFs

- Main hyperparameter: `max_features`
 - Number of features for each split
 - \sqrt{p} is common for classification
 - $\sim p$ for regression
- Number of estimators
 - More is better
- Pre-pruning can help save memory
 - `Max_depth`, `max_leaf_nodes`,
`min_samples_split`, etc.

Stacked Generalization - Stacking

- Combine models in a different way (meta-learners)
 - combine learners of different types
- Idea
 - split training data into two sets
 - train several learners on first part
 - test these learners on second part
 - use the test predictions as inputs and target output to train a higher level of learner

Strong vs. Weak Learners

- So far, strategy has been:
 - gather a bunch of data
 - think hard, then make a single, large, complicated predictor
 - test the predictor on data
- What if we wanted to use a bunch of simple predictors instead, is there a principled way to do this?

Strong vs. Weak Learners

- A strong learner is a method that can learn a decision rule arbitrarily well
- A weak learner is a simple method that does better than guessing, but cannot learn a decision rule arbitrarily well.
- Example: trying to decide whether email is spam
 - Strong learner: method that uses words, syntax, etc. as features, and fits a high-accuracy decision rule
 - Weak learner: use simple rules, if phrase "deal available" is in email, then predict is spam

Boosting

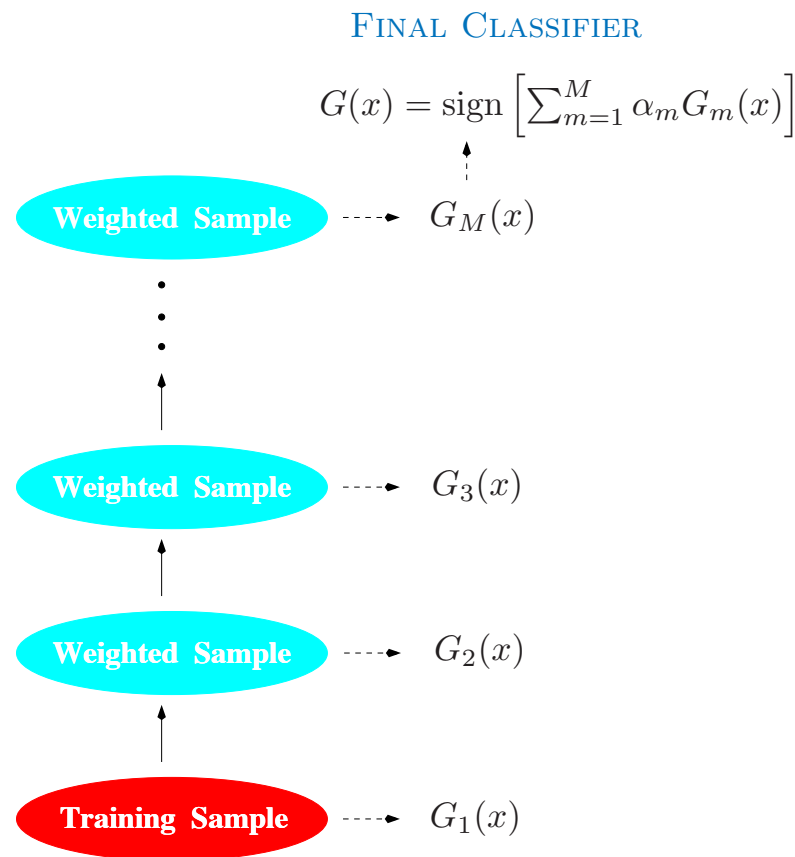
- Powerful technique originally designed for classification
 - extended to regression
- Basic Idea:
 - use a “weak” classifier (accuracy only slightly better than random)
 - create a series of such classifiers where the training data that was mis-classified on the previous iteration is given additional weight
 - combine successive models by voting to create a final model
- Example: Adaptive Boosting (AdaBoost)

Boosting

- Learning over weighted training set
 - **Weights** are assigned to each training tuple
 - A series of k classifiers is iteratively learned
 - After a classifier M_i is learned, the weights are updated to allow the subsequent classifier, M_{i+1} , to **pay more attention to the training tuples that were misclassified by M_i**
 - The final **M^* combines the votes** of each individual classifier, where the weight of each classifier's vote is a function of its accuracy

Boosting

- Difference between bagging and boosting
 - In boosting fit model to the entire training set, but adaptively weight the samples



AdaBoost (Freund and Schapire, 1997)

- Given data set, $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$
- $W(x)$ is a distribution of weights over the n training examples
- Initially, set uniform weight distribution $w_i = 1/n$
- For each iteration, k
 - find model (hypothesis) $H_k(x)$ with min. error e_h using weights $W_k(x)$
 - compute α_k
 - Update weights
 - correctly labeled samples; decrease wt $W_{k+1} = W_k * \exp(-\alpha_k)$
 - incorrectly labeled samples; increase wt $W_{k+1} = W_k * \exp(\alpha_k)$
- Final Model

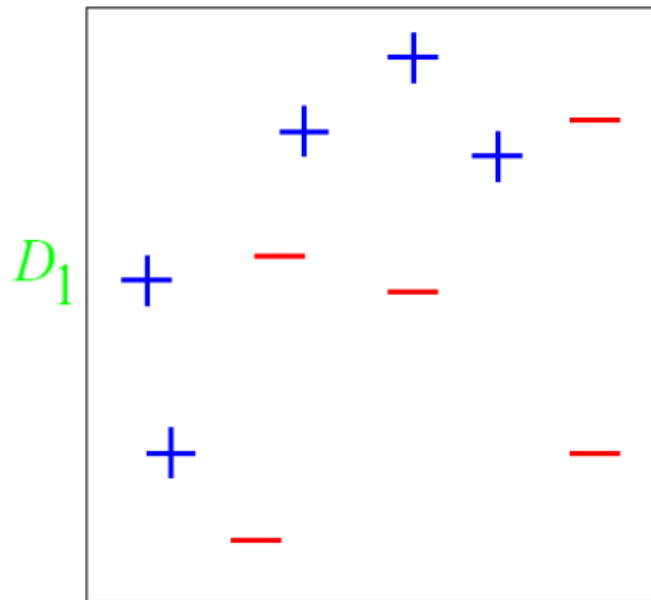
$$H_{final}(x) = \text{sign}(\sum \alpha_k H_k(x))$$

AdaBoost

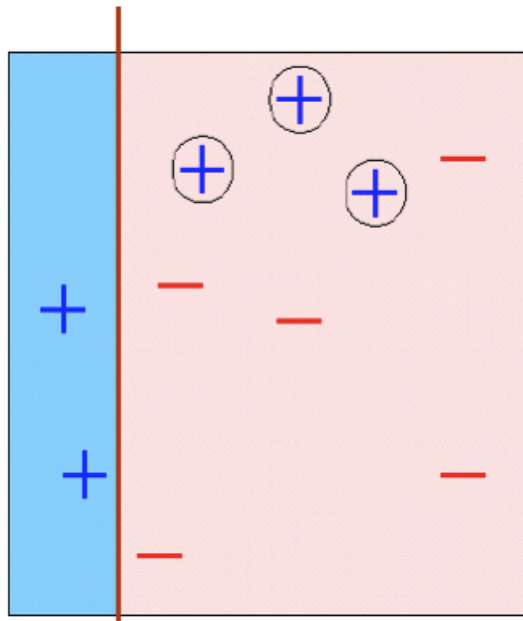
- Why consider using AdaBoost?
 - No tunable parameters
 - Works with any weak learner
 - Computational feasible
 - Tends to avoid overfitting

Example: AdaBoost

- Initial Training data
- All weights equal, for each sample



Example: AdaBoost



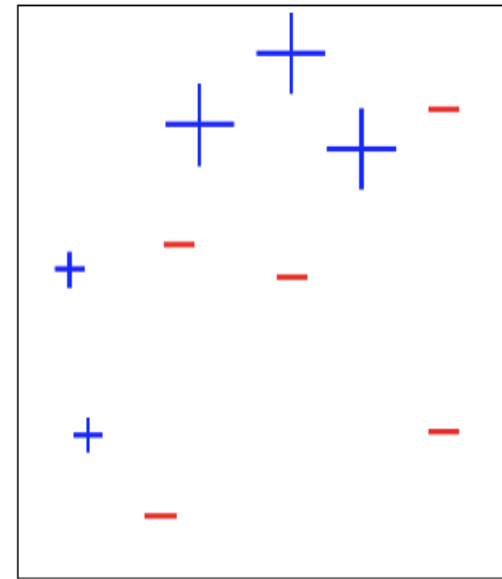
h_1

$$e_1 = 0.3$$

$$\alpha_1 = 0.42$$



D_2



$$e_1 = \sum_{\text{incorrect}} w(\text{incorrect}) = 0.3$$

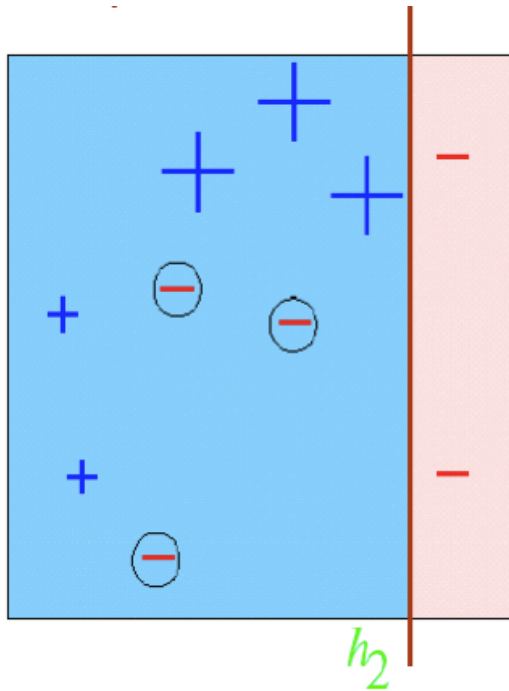
$$\alpha_1 = \frac{1}{2} \ln \left(\frac{1 - e_1}{e_1} \right) = \frac{1}{2} \ln \left(\frac{0.7}{0.3} \right) = 0.42$$

$$w_{\text{corr}}' = C_N w_{\text{corr}} e^{-\alpha_1} = (1.091) 0.1 (0.6546) = 0.0714$$

$$w_{\text{incorr}}' = C_N w_{\text{incorr}} e^{\alpha_1} = (1.091) 0.1 (1.5275) = 0.1667$$

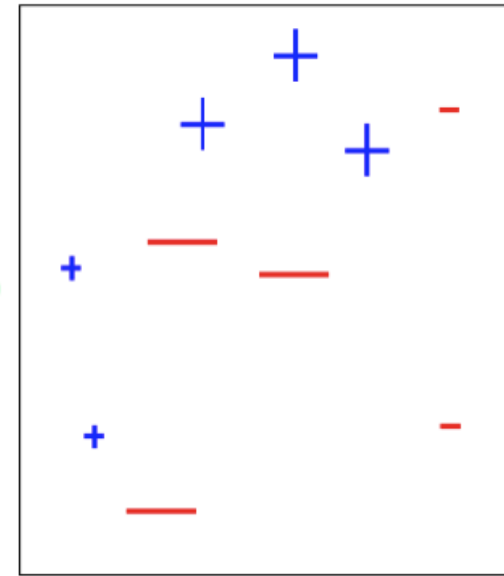
$$C_N = 1 / \sum_N w' = 1.091$$

Example: AdaBoost



$$e_2 = 0.21$$

$$\alpha_2 = 0.65$$



$$e_2 = \sum_{\text{incorrect}} w(\text{incorrect}) = 0.21$$

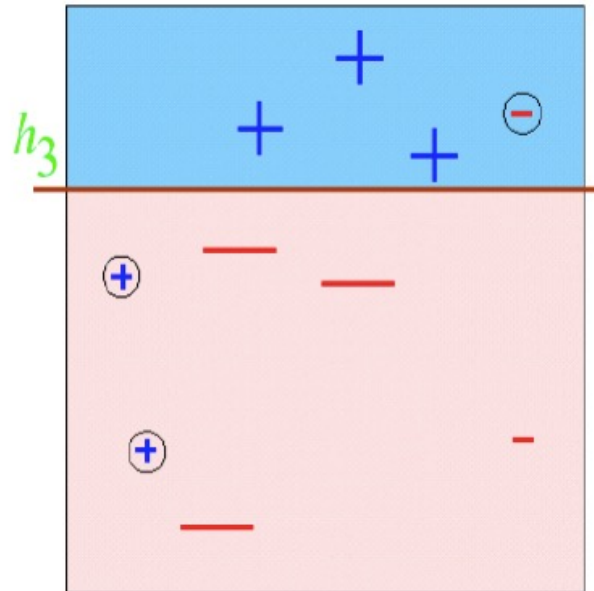
$$\alpha_2 = \frac{1}{2} \ln \left(\frac{1 - e_1}{e_1} \right) = \frac{1}{2} \ln \left(\frac{0.79}{0.21} \right) = 0.65$$

$$w_{\text{corr}}' = C_N w_{\text{corr}} e^{-\alpha_2}$$

$$w_{\text{incorr}}' = C_N w_{\text{incorr}} e^{\alpha_2}$$

$$C_N = 1 / \sum_N w'$$

Example: AdaBoost

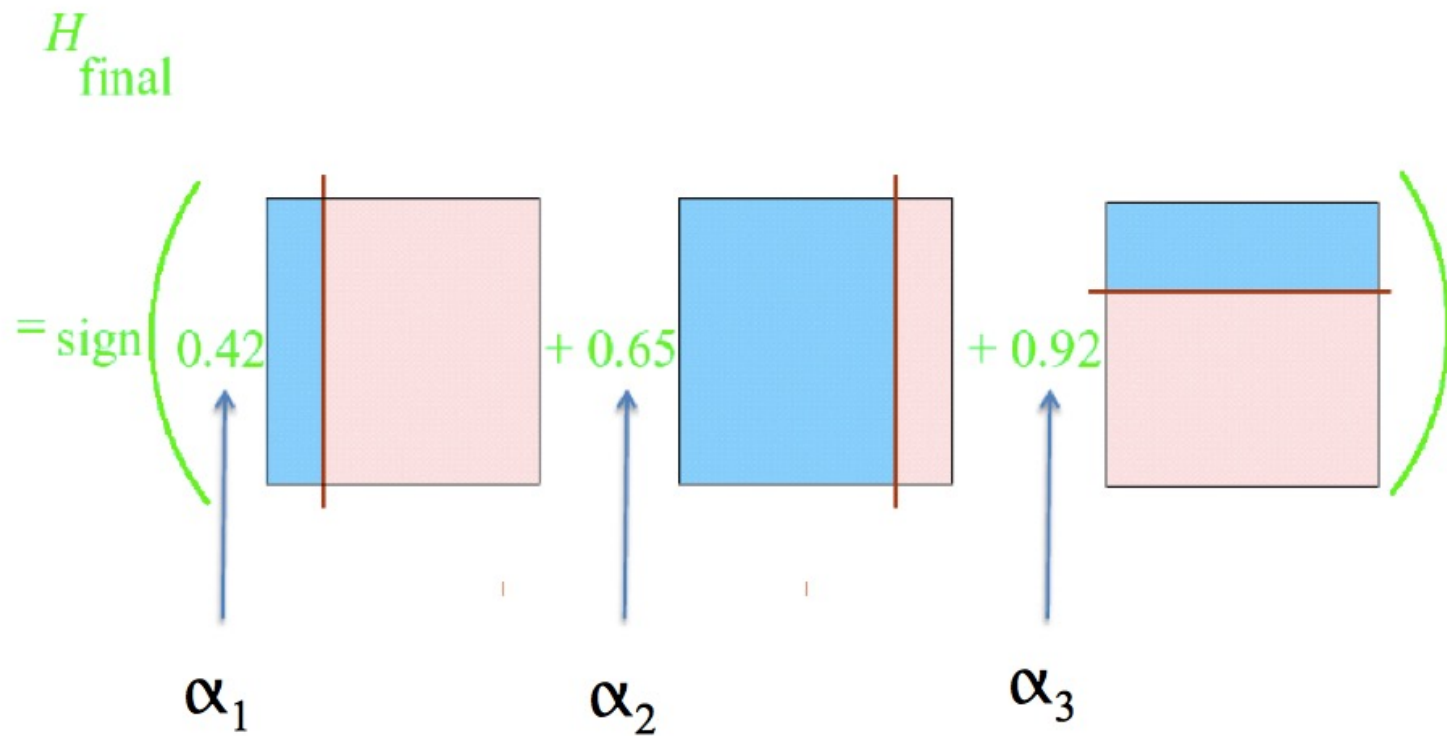


$$e_3 = 0.14$$

$$\alpha_3 = 0.92$$

Example: AdaBoost

Final Model (Hypothesis):

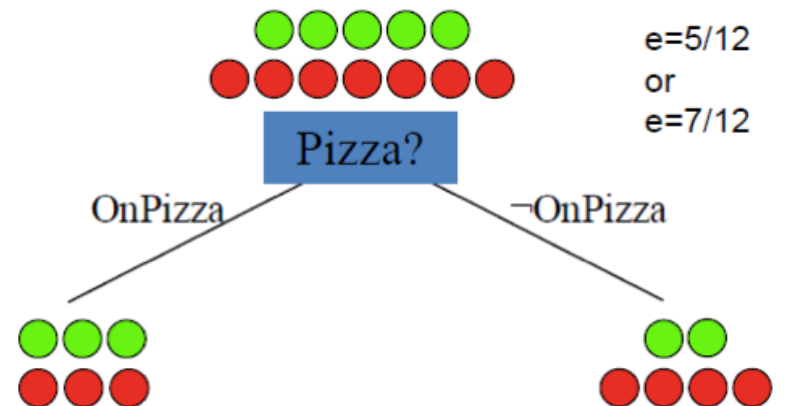
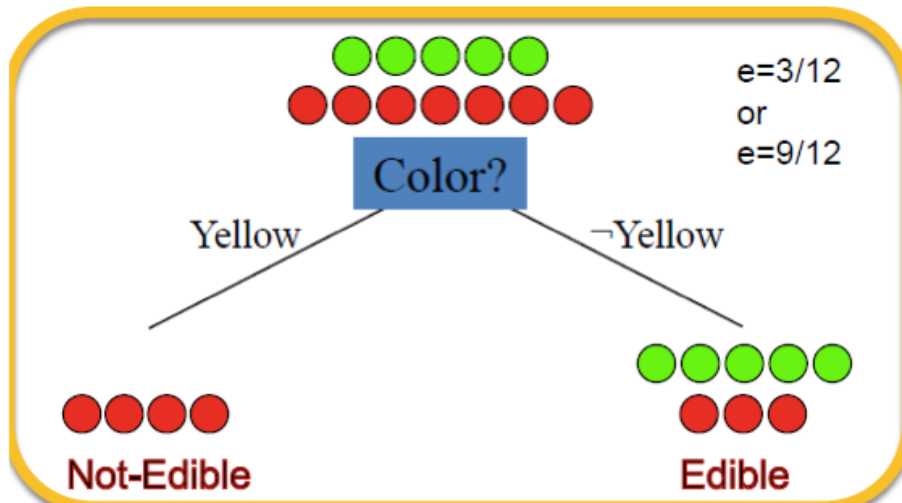
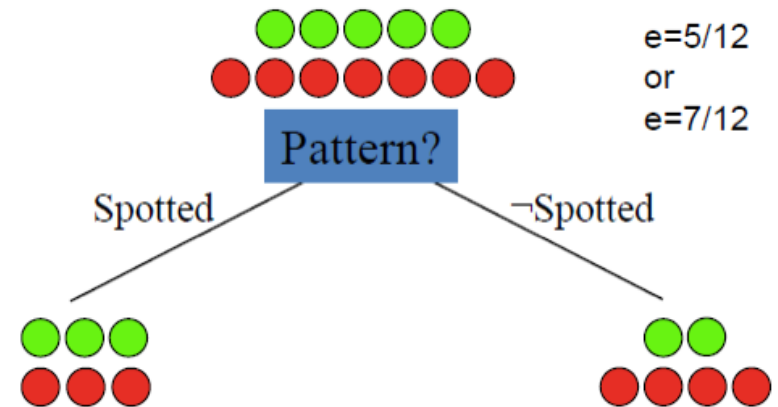
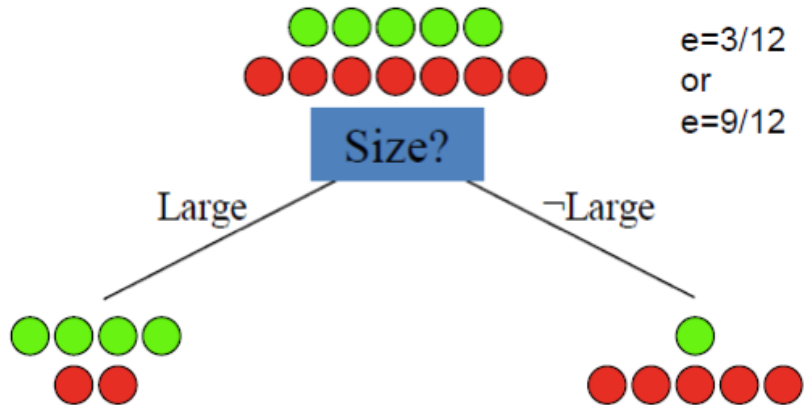


Example: AdaBoost Mushroom

- Mushroom data set used
- Initial weights = $1/12 = 0.0833$

Pattern	Size	Color	OnPizza	Edible
S	L	Y	Y	No
S	L	N	Y	Yes
S	L	N	N	Yes
S	S	Y	N	No
S	S	N	Y	Yes
S	S	N	N	No
N	L	Y	N	No
N	L	N	Y	Yes
N	L	N	N	Yes
N	S	Y	Y	No
N	S	N	Y	No
N	S	N	N	No

Boosting on Features



Boosting on Features

- Compute Weights
- Decision stump is on Color:
 - Yellow = not edible, not yellow = edible

$$e_1 = \sum w(\text{incorrect}) = 1/12 + 1/12 + 1/12 = 1/4$$

$$\alpha_1 = \frac{1}{2} \ln \frac{3}{1} = 0.55$$

$$w(\text{correct}) = w(\text{incorrect}) = 1/12 = 0.0833$$

$$w'(\text{correct}) = 0.0555$$

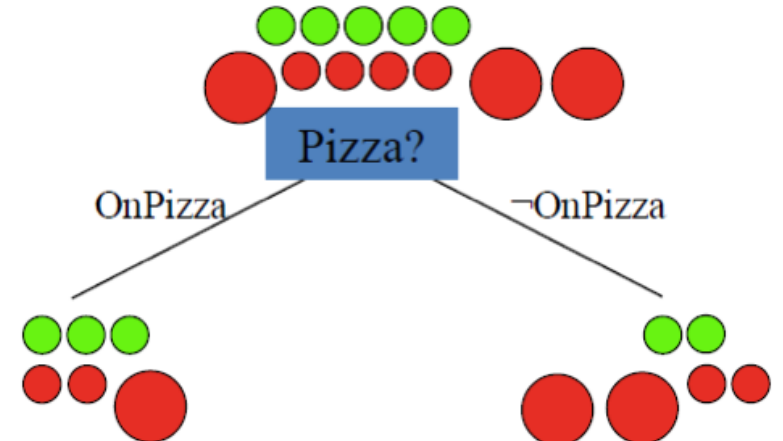
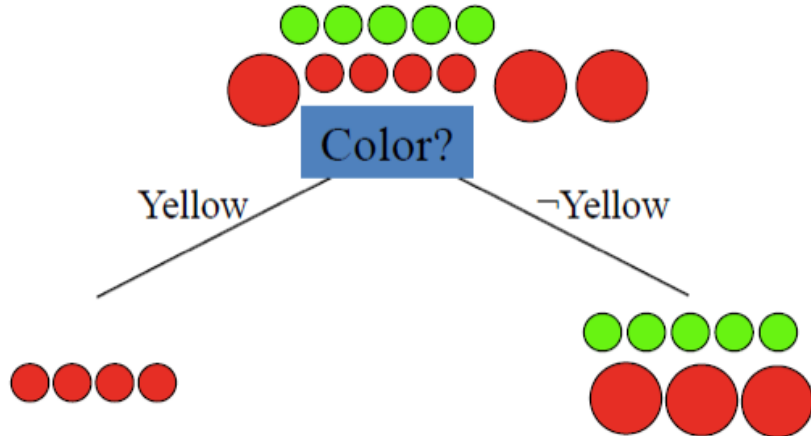
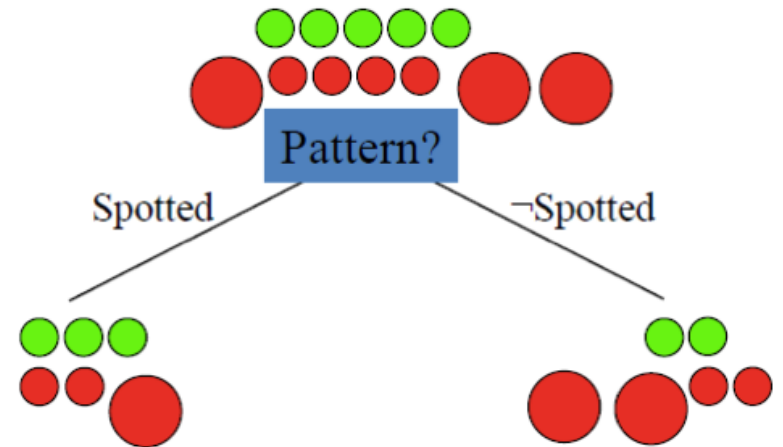
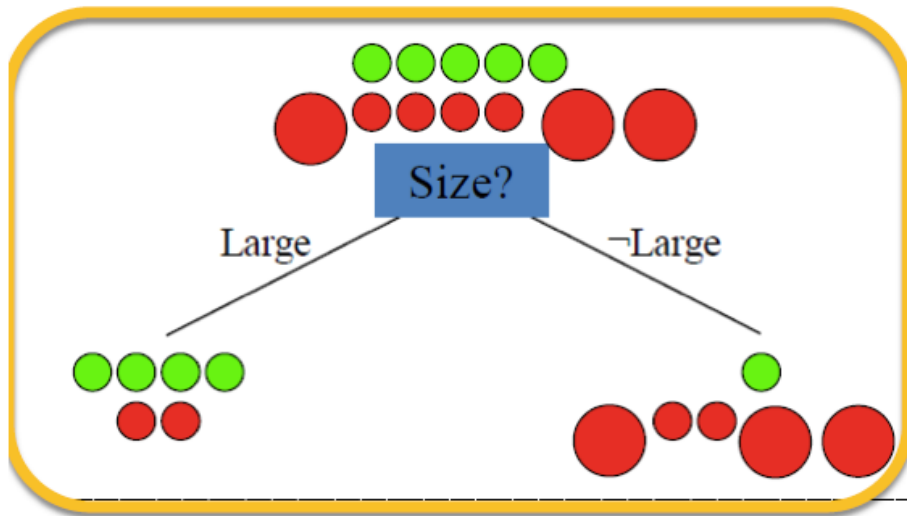
$$w'(\text{incorrect}) = 0.1666$$

Example: Next Iteration

- Weighted data

Pattern	Size	Color	OnPizza	Edible	
S	L	Y	Y	No	.0555
S	L	N	Y	Yes	.0555
S	L	N	N	Yes	.0555
S	S	Y	N	No	.0555
S	S	N	Y	Yes	.0555
S	S	N	N	No	.1666
N	L	Y	N	No	.0555
N	L	N	Y	Yes	.0555
N	L	N	N	Yes	.0555
N	S	Y	Y	No	.0555
N	S	N	Y	No	.1666
N	S	N	N	No	.1666

Boosting on Features (2)



Boosting on Features (2)

- Compute Weights
- Decision stump is for size
 - Large = Edible, not large = not edible

$$e_2 = \sum w(\text{incorrect}) = (2 * 0.0555) + (1 * 0.0555) = 0.1665$$

$$\alpha_2 = \frac{1}{2} \ln\left(\frac{0.8335}{0.1665}\right) = 0.80$$

Example: Boosting

- ESL p. 339, data with $n=1000$ points
- A single stump produces a misclassification rate of 45.8%
- With boosting with 400 iterations, the misclassification rate is 5.8%
- This also beats the misclassification rate of a single tree – 24.7%

Example: Boosting

