

Text Mining and Information Retrieval

Laura Brown

Some slides adapted from P. Smyth; Han, Kamber, & Pei;
Tan, Steinbach, & Kumar; C. Volinsky; R. Tibshirani; D. Kauchak
and <http://nlp.stanford.edu/IR-book/>

Outline

- Information Retrieval
 - What is it?
 - Challenges
- Represent Text Data
 - Term-document incidence matrices
 - Inverted index
- Boolean Queries
 - Query processing
 - Query optimization

Information Retrieval

- What comes to mind when I say “information retrieval”?
- Where have you seen IR? What are some real-world examples/uses?
 - Search engines (web search)
 - File search (e.g. OS X Spotlight, Windows Instant Search, Google Desktop)
 - Databases?
 - Catalog search (e.g. library)
 - Intranet search (i.e. corporate networks)

Information Retrieval

- Information Retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers).

Information Retrieval

- Information Retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers).
 - Find all documents about computer science
 - Find all course web pages at Michigan Tech
 - What is the cheapest flight from LA to NY?
 - Who was the 15th president?

Information Retrieval

- Information Retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers).
- What is the difference between an *information need* and a *query*?

Information Retrieval

- Information Retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers).

Information need

- Find all documents about computer science
- Find all course web pages at Michigan Tech
- Who is was the 15th president?

Query

“computer science”

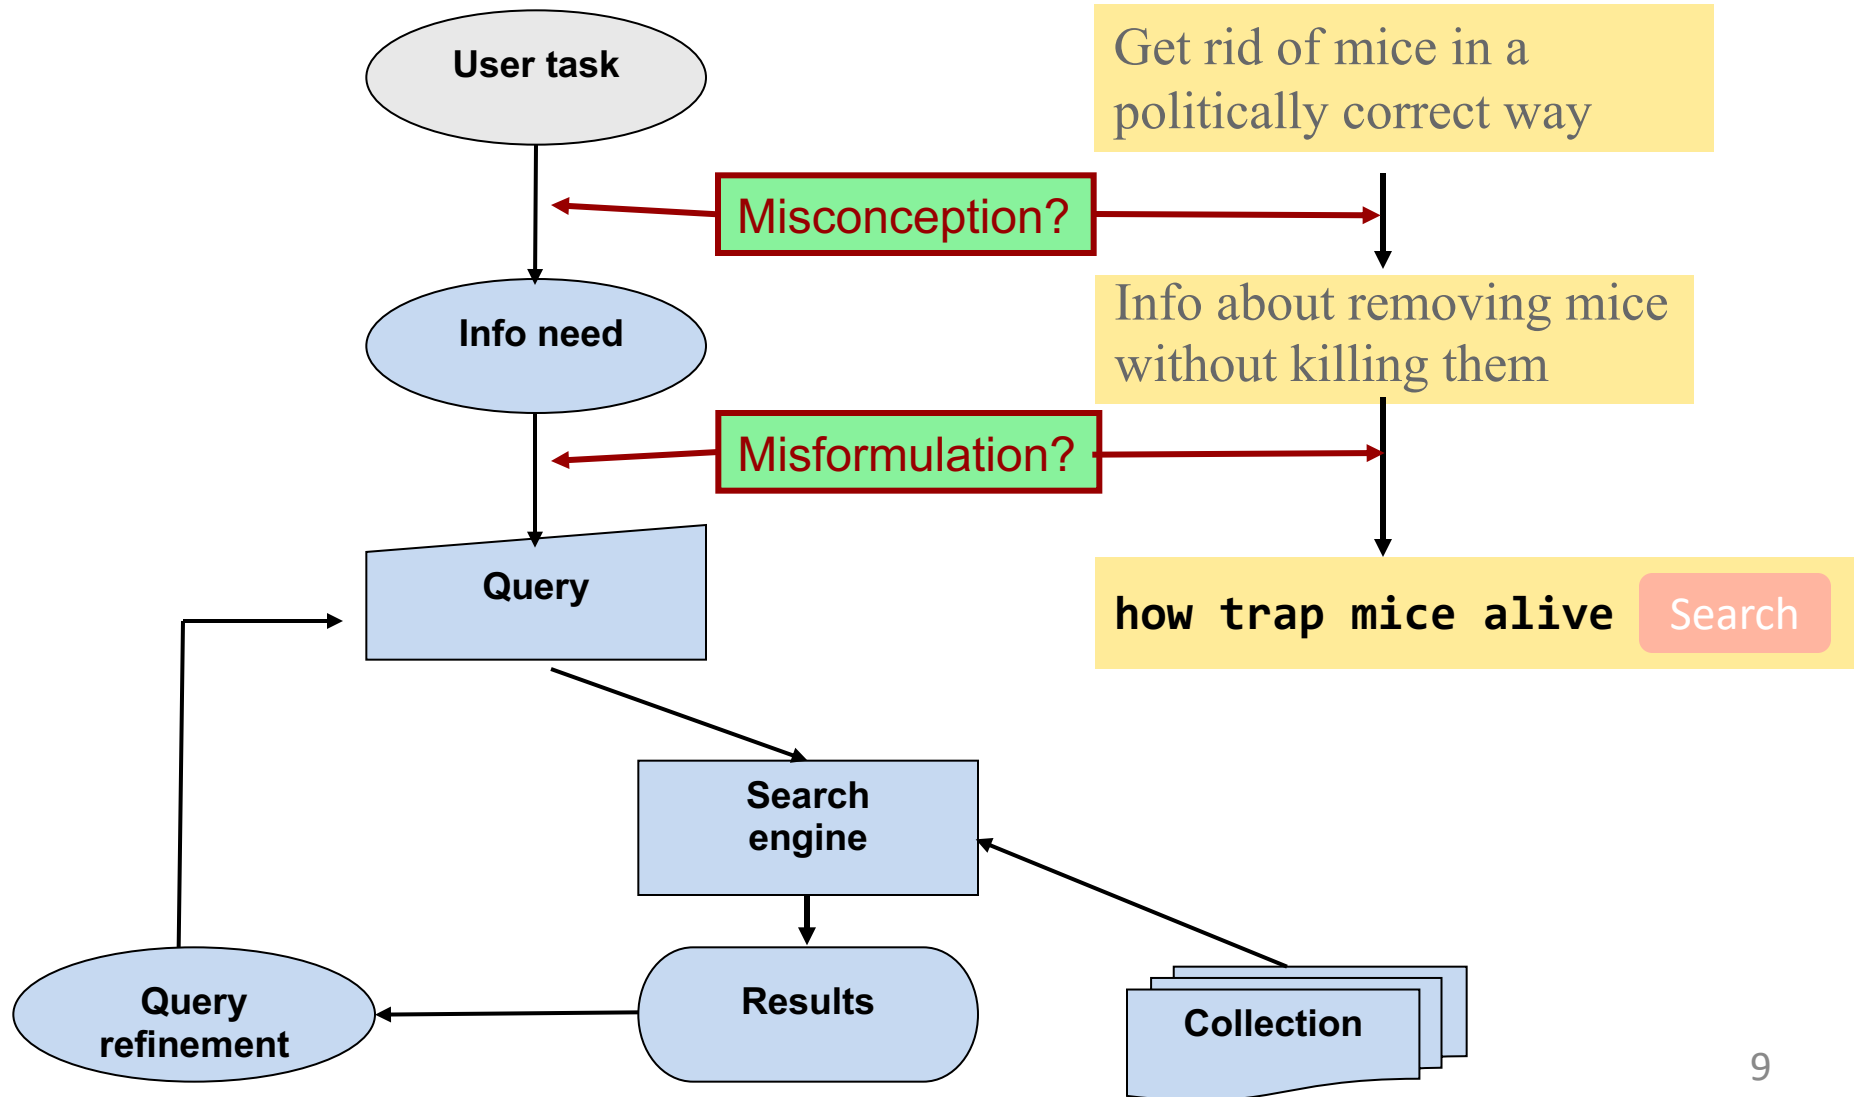
Michigan Tech AND college
AND *url-contains* class

WHO=president NUMBER=15

Assumptions of Information Retrieval

- **Collection**: A set of documents
 - Assume it is a static collection for the moment
- **Goal**: Retrieve documents with information that is **relevant** to the user's **information need** and helps the user complete a **task**

Classic Search Model



How good are the retrieved docs?

- *Precision* : Fraction of retrieved docs that are relevant to the user's **information need**
- *Recall* : Fraction of relevant docs in collection that are retrieved
 - More precise definitions and measurements to follow later

Information Retrieval Challenges

- Main problems with text retrieval:
 - What does relevant mean?
 - How do you know if you have the right documents?
 - How can user feedback be incorporated?

Term-document incidence matrices

Introduction to Information Retrieval

Unstructured data in 1620

- Which plays of Shakespeare contain the words ***Brutus AND Caesar*** but ***NOT Calpurnia***?
- One could **grep** all of Shakespeare's plays for ***Brutus*** and ***Caesar***, then strip out lines containing ***Calpurnia***?
- Why is this not the answer?

Unstructured data in 1620

- Which plays of Shakespeare contain the words ***Brutus AND Caesar*** but ***NOT Calpurnia***?
- One could `grep` all of Shakespeare's plays for ***Brutus*** and ***Caesar***, then strip out lines containing ***Calpurnia***?
- Why is this not the answer?
 - Slow (for large corpora)
 - ***NOT Calpurnia*** is non-trivial
 - Other operations (e.g., find the word ***Romans*** near ***countrymen***) not feasible
 - Ranked retrieval (best documents to return)

Term-document incidence matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Brutus AND Caesar BUT NOT Calpurnia

1 if play contains
word, 0 otherwise

Incidence vectors

- For each term, we have a 0/1 vector
 - Caesar = 110111
 - Brutus = 110100
 - Calpurnia = 010000
- To answer query?

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Incidence vectors

- For each term, we have a 0/1 vector
 - Caesar = 110111
 - Brutus = 110100
 - Calpurnia = 010000
- To answer query? take vectors for **Brutus, Caesar** and **Calpurnia** (complemented) → bitwise *AND*.
 - Answer = 100100

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0


Answers to query

- Antony and Cleopatra, Act II, Scene ii

Agrippa [Aside to DOMITIUS ENOBARNUS]: Why, Enobarbas,
When Antony found Julius **Caesar** dead,
He cried almost to roaring; and he wept
When at Philippi he found **Brutus** slain.

- Hamlet, Act III, Scene ii

Lord Polonius: I did enact Julius **Caesar** I was killed I' the
Capitol; **Brutus** killed me.



Incidence vectors

- For each term, we have a 0/1 vector
 - Caesar = 110111
 - Brutus = 110100
 - Calpurnia = 010000
- Bitwise AND the vectors together using the complemented vector for all NOT queries

Any problem with this approach?

Bigger Collections

- Consider $N = 1$ million documents, each with about 1000 words.
- Ave. 6 bytes/word including spaces/punctuation
 - 6 GB of data in the documents
- Say there are $M = 500\text{K}$ *distinct* terms among these.

Bigger Collections



- Can't build the matrix!
- 500K x 1M matrix has half-a-trillion 0's and 1's.
- But it has no more than one billion 1's.
 - Each of the 1 million documents has at most 1000 1's
 - Matrix is extremely sparse!
- What's a better representation?
 - Only record the 1 positions

Inverted Index


Introduction to Information Retrieval

Inverted index

- For each term t , we store a list of all documents that contain t
 - Identify each doc by a **docID**, a document serial number
- What data structures might we use for this?




Inverted index representation




- Pros
 - Simple to implement
 - No extra pointers required for data structure
 - Contiguous memory
- Cons
 - How do we pick the size of the array?
 - What if we want to add additional documents?

Inverted index representation



- Pros
 - Dynamic space allocation
 - Insertion of new documents is straightforward
- Cons
 - Memory overhead of pointers
 - Noncontiguous memory access

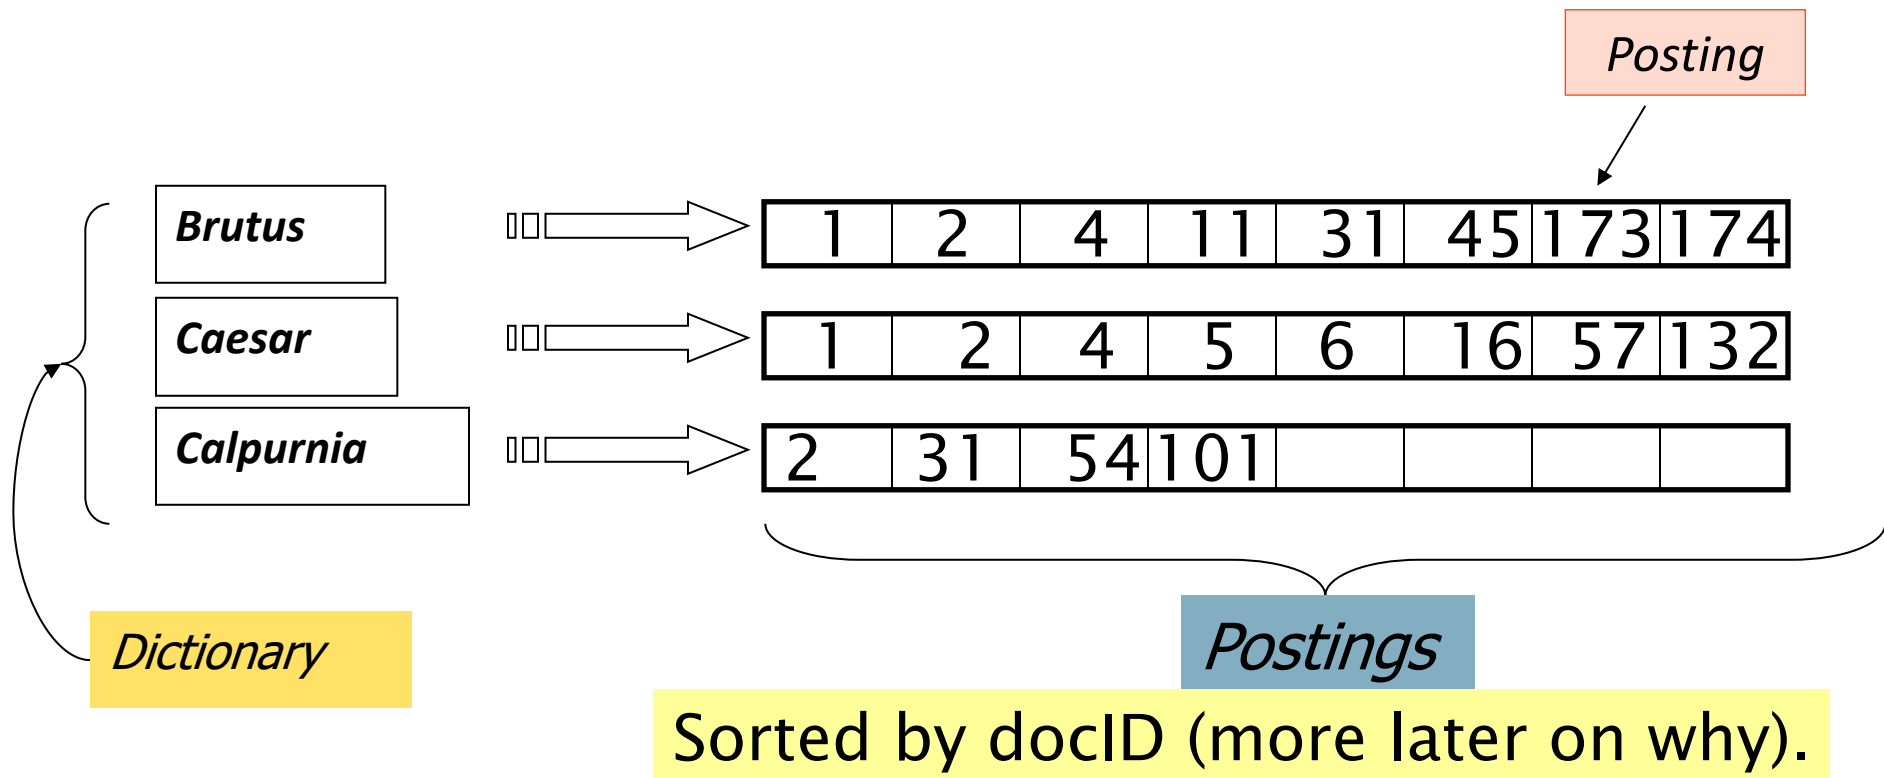
Inverted index representation



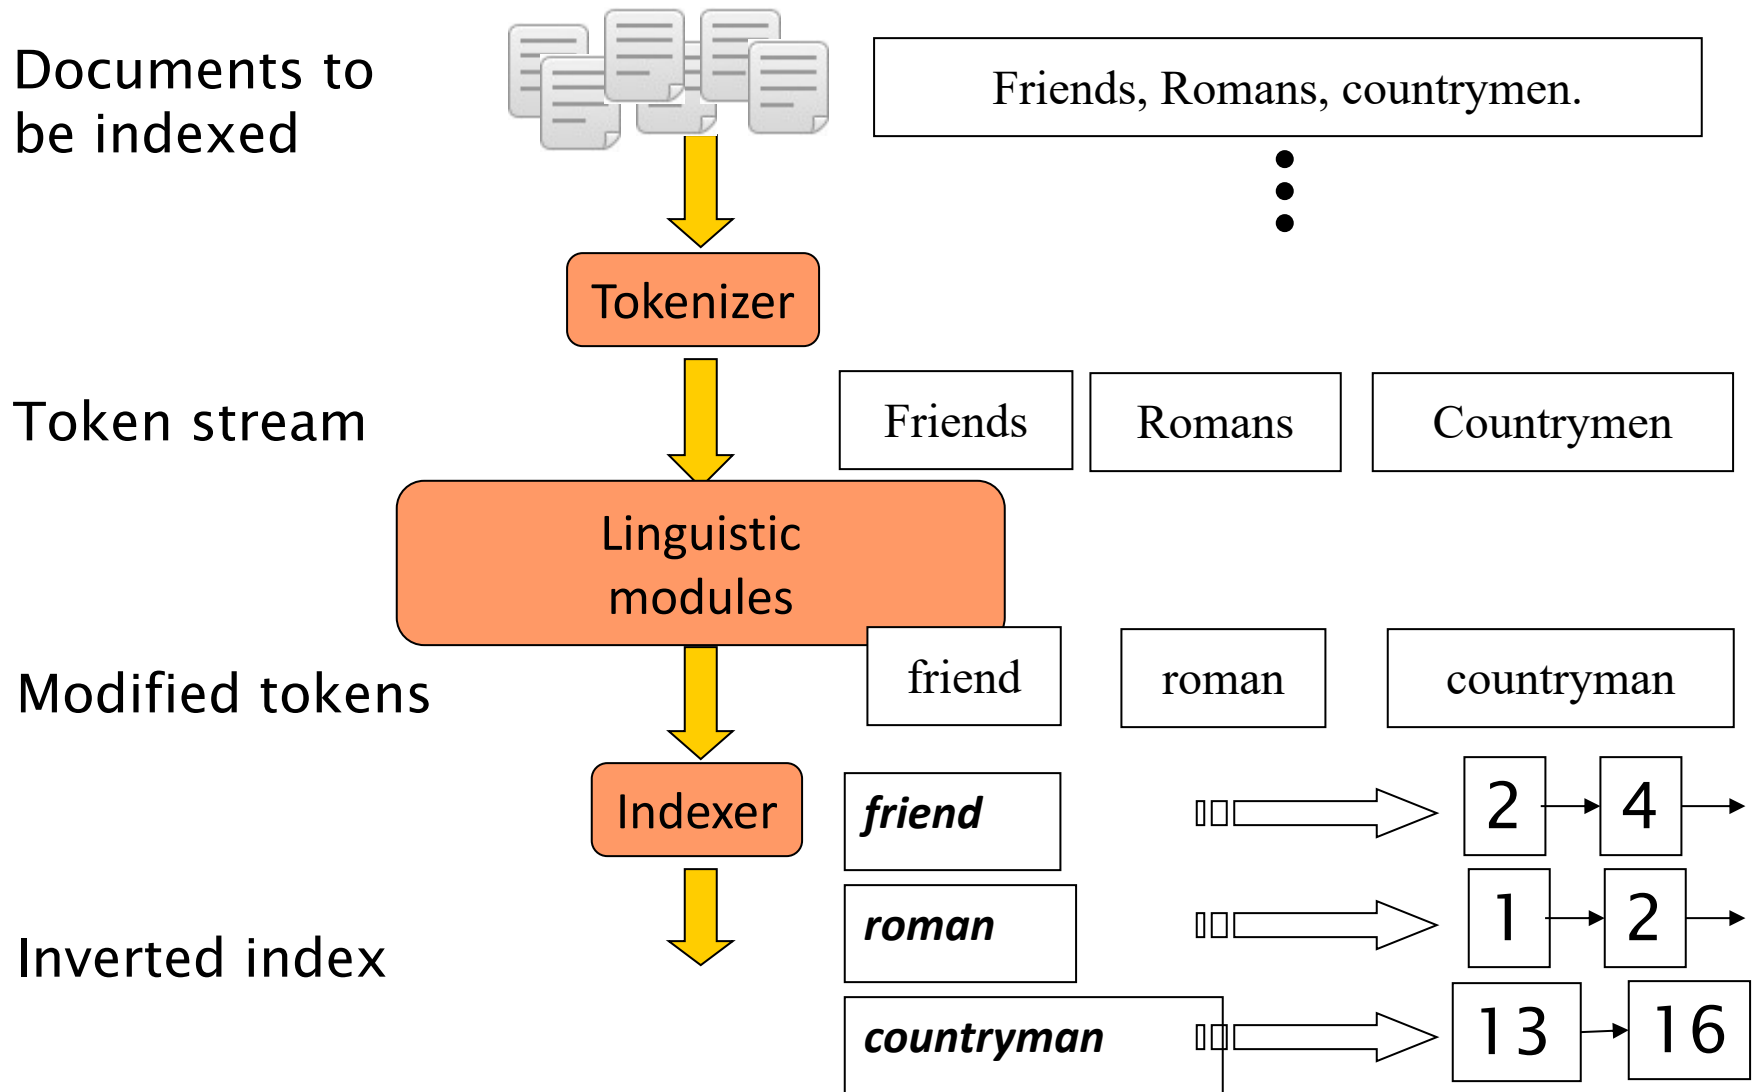
- Pros
 - Search in constant time
 - Contiguous memory
- Cons
 - How do we pick the size?
 - What if we want to add additional documents?
 - May have to rehash
 - To get constant time operations, lots of unused slots/memory

Inverted index

- We need variable-size **postings lists**
 - On disk, a continuous run of postings
 - In memory, can use linked lists



Inverted index construction



Initial stages of text processing

- Tokenization

- Cut character sequence into word tokens
 - Deal with ***“John’s”, a state-of-the-art solution***

- Normalization

- Map text and query term to same form
 - You want **U.S.A.** and **USA** to match

- Stemming

- We may wish different forms of a root to match
 - ***authorize, authorization***

- Stop words

- We may omit very common words (or not)
 - ***the, a, to, of***

We will dig into these later!

Indexer steps: Token sequence

- Sequence of (Modified token, Document ID) pairs.

Doc 1

I did enact Julius
Caesar I was killed
i' the Capitol;
Brutus killed me.

Doc 2

So let it be with
Caesar. The noble
Brutus hath told you
Caesar was ambitious



Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2

Indexer steps: Sort

- Sort by terms
 - At least conceptually
 - And then docID



Core indexing step

Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2



Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
I	1
I	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2

Indexer steps: Dictionary & Postings

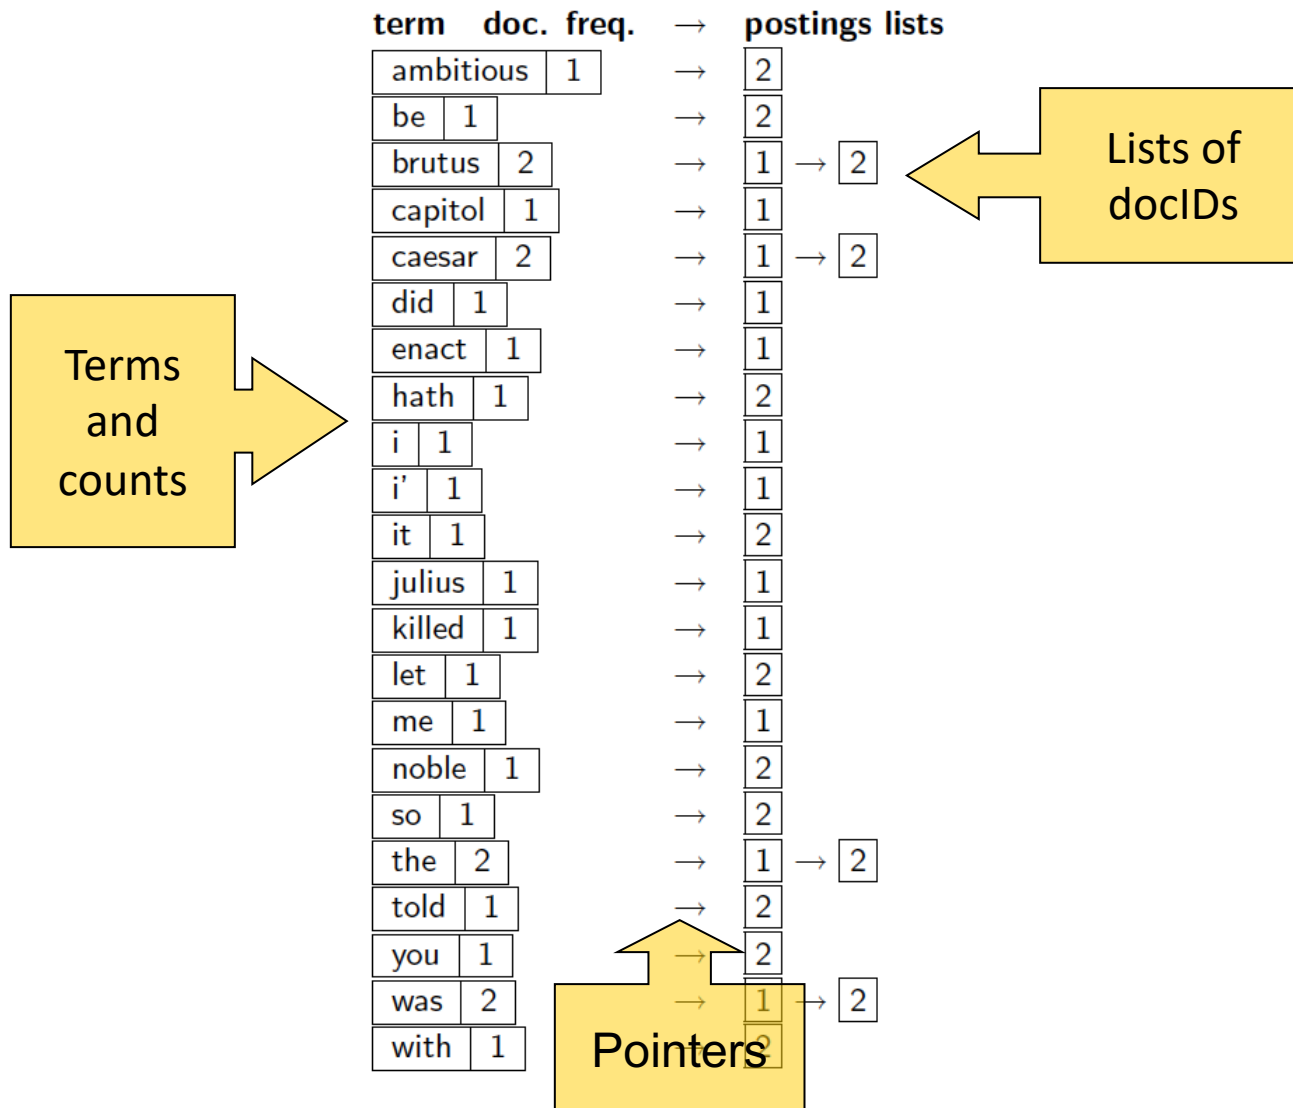
- Multiple term entries in a single document are merged.
- Split into Dictionary and Postings
- Doc. frequency information is added.

Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
i	1
i	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2



term	doc. freq.	→	postings lists
ambitious	1	→	2
be	1	→	2
brutus	2	→	1 → 2
capitol	1	→	1
caesar	2	→	1 → 2
did	1	→	1
enact	1	→	1
hath	1	→	2
i	1	→	1
i'	1	→	1
it	1	→	2
julius	1	→	1
killed	1	→	1
let	1	→	2
me	1	→	1
noble	1	→	2
so	1	→	2
the	2	→	1 → 2
told	1	→	2
you	1	→	2
was	2	→	1 → 2
with	1	→	2

Indexing Storage



Query processing with an inverted index


Introduction to Information Retrieval

Query processing: AND

- Consider processing the query:



Brutus AND Caesar

- Locate ***Brutus*** in the Dictionary;
 - Retrieve its postings
- Locate ***Caesar*** in the Dictionary;
 - Retrieve its postings
- “Merge” the two postings (intersect the document sets)



The merge



- Walk through the two postings simultaneously



Brutus AND Caesar

The merge



- Walk through the two postings simultaneously




Brutus AND Caesar

The merge

- Walk through the two postings simultaneously





Brutus AND Caesar




The merge

- Walk through the two postings simultaneously





Brutus AND Caesar




The merge

- Walk through the two postings simultaneously





Brutus AND Caesar




The merge

- Walk through the two postings simultaneously





Brutus AND Caesar



The merge

- Walk through the two postings simultaneously





...



The merge

- Walk through the two postings simultaneously





What assumption are we making about the postings lists?

For efficiency, when we construct the index, we ensure that the postings lists are sorted

The merge

- Walk through the two postings simultaneously



What is the running time?

$O(\text{length1} + \text{length2})$

Boolean Queries

Introduction to Information Retrieval

Boolean queries: exact match

- The **Boolean retrieval model** is being able to ask a query that is a Boolean expression:
 - Boolean Queries are queries using *AND*, *OR* and *NOT* to join query terms
 - Views each document as a set of words
 - Is precise: document matches condition or not.
 - Perhaps the simplest model to build an IR system on
- Primary commercial retrieval tool for 3 decades.
- Many search systems you still use are Boolean:
 - Email, library catalog, macOS Spotlight

Merging

What about an arbitrary Boolean formula?

(Brutus OR Caesar) AND NOT


(Antony OR Cleopatra)

- $x = (\text{Brutus OR Caesar})$
 - $y = (\text{Antony OR Cleopatra})$
 - $x \text{ AND NOT } y$
- Is there an upper bound on the running time?
 - $O(\text{total_terms} * \text{query_terms})$
 - What about **Brutus AND Calpurnia AND Caesar?**

Query Optimization

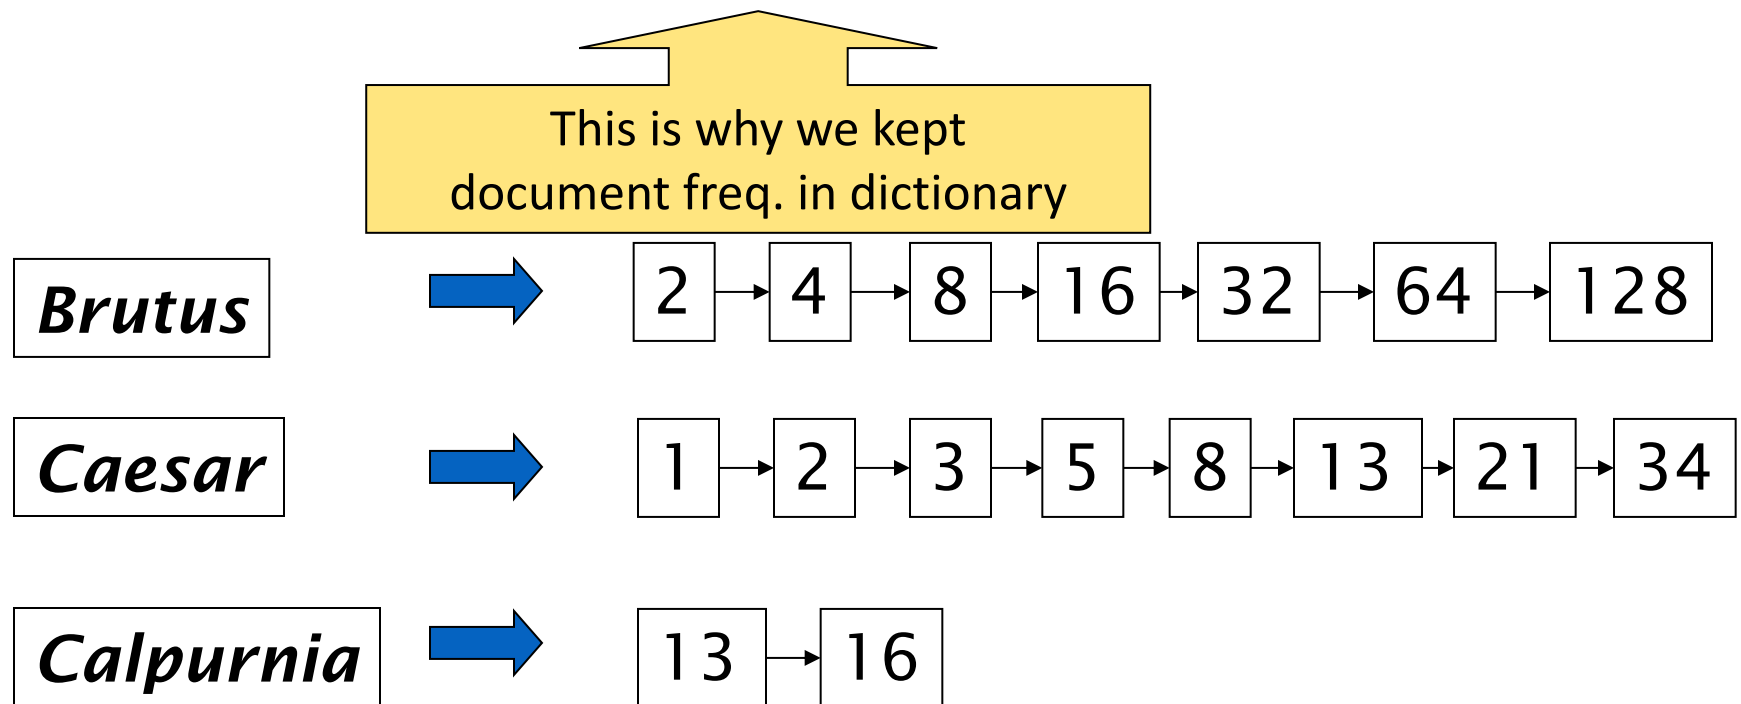
Query: *Brutus AND Calpurnia AND Caesar*

- Consider a query that is an *AND* of t terms.
- For each of the terms, get its postings, then *AND* them together
- What is the best order for query processing?



Query optimization example

- Heuristic: **Process in order of increasing freq:**
 - *merge the two terms with the shortest postings list*




Execute the query as (**Calpurnia** **AND** **Brutus**) **AND** **Caesar**.

Query optimization - OR

Query: *Brutus OR Calpurnia OR Caesar*

- Consider a query that is an *OR* of t terms.
- What is the best order for query processing?
- Same: still want to merge the shortest postings lists first



Query optimization in general

Query: (madding OR crowd) AND (ignoble OR strife)

- Need to evaluate OR statements first
- Which OR should we do first?
 - Get doc. freq.'s for all terms
 - Estimate the size of each OR by the sum of its doc. freq.'s (conservative)
 - Process in increasing order of OR sizes